

Machine Learning Group Assignment

March 28, 2025

E-Portfolio on the UK Road Safety Dataset: Group Assignment



Candidate Numbers: 141715, 835522, 750476, 788719

Table of Contents

1. Importing Libraries and Preparing Environment
2. Business Objective
3. Data Preparation
4. Data Splitting
5. Exploratory Data Analysis
6. Data Preprocessing
7. Conclusion
8. Data Exporting

1 Importing Libraries and Preparing Environment

```
[2]: #Base Libraries
import numpy as np
import pandas as pd
import warnings
```

```

from scipy import stats
from timeit import default_timer as timer
from datetime import timedelta

#Library for Plotting
import seaborn as sns
import matplotlib.pyplot as plt

#Library for Data Preprocessing and Cleaning
from sklearn.exceptions import ConvergenceWarning
from sklearn.model_selection import train_test_split, StratifiedShuffleSplit
from sklearn.base import TransformerMixin, BaseEstimator
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer, SimpleImputer
from sklearn.preprocessing import OneHotEncoder
from sklearn.ensemble import IsolationForest

#Prevent warnings of the code from showing
warnings.filterwarnings(action = 'ignore')

```

2 Context and Project Objective

In 2023, approximately 23,000 road collisions occurred across London boroughs; while most caused only minor injuries, around 3,500 involved burns or fractures, and 101 resulted in at least one fatality (Trust For London, 2025). Four datasets from the Department for Transport were merged, cleaned, and split into training and testing sets. Irrelevant variables were removed, and the data underwent imputation, log transformation, and outlier detection.

This project aims to build a machine learning model to predict the severity of traffic accidents using variables related to driver behavior, vehicle type, and environmental conditions. The model will support Lambeth, Wandsworth, Croydon, Westminster & Southwark boroughs to enhance road safety through targeted interventions and awareness campaigns.

3 Data Preparation

3.1 Loading Data

```

[7]: #Load in the datasets
vehicle_data = pd.read_csv('dft-road-casualty-statistics-vehicle-2023.csv')
casualty_data = pd.read_csv('dft-road-casualty-statistics-casualty-2023.csv')
collision_data = pd.read_csv('dft-road-casualty-statistics-collision-2023.csv')
guide_data = pd.
    ↪read_excel('dft-road-casualty-statistics-road-safety-open-dataset-data-guide-2024.
    ↪xlsx')

```

```

[8]: #Verify the data is loaded in correctly
vehicle_data.head()

```

```

[8]:  accident_index  accident_year  accident_reference  vehicle_reference  \
0  2023010419171      2023      10419171      1
1  2023010419183      2023      10419183      1
2  2023010419183      2023      10419183      2
3  2023010419183      2023      10419183      3
4  2023010419189      2023      10419189      1

    vehicle_type  towing_and_articulation  vehicle_manoeuvre  \
0           11              0              4
1           11              0             18
2           9              0              9
3           9              0              8
4           9              0             18

    vehicle_direction_from  vehicle_direction_to  \
0              1              5
1              5              1
2              1              6
3              7              1
4              7              3

    vehicle_location_restricted_lane  ...      generic_make_model  \
0              0  ...  ALEXANDER DENNIS MODEL MISSING
1              0  ...      WRIGHTBUS GEMINI
2              0  ...      TOYOTA YARIS
3              0  ...      BMW 2 SERIES
4              0  ...      LEXUS RX 400

    driver_imd_decile  driver_home_area_type  lsoa_of_driver  scooter_flag  \
0              3              1      E01001177      0
1              6              1      E01001419      0
2              3              1      E01001546      0
3              4              1      E01001686      0
4              5              1      E01002443      0

    dir_from_e  dir_from_n  dir_to_e  dir_to_n  driver_distance_banding
0      NaN      NaN      NaN      NaN      2
1      NaN      NaN      NaN      NaN      2
2      NaN      NaN      NaN      NaN      1
3      NaN      NaN      NaN      NaN      4
4      NaN      NaN      NaN      NaN      1

[5 rows x 34 columns]

```

```

[9]:  #Verify the total number of samples and features of the data
      vehicle_data.shape

```

[9]: (189815, 34)

```
[10]: #Check the information of the dataset
vehicle_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 189815 entries, 0 to 189814
Data columns (total 34 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   accident_index                        189815 non-null  object
1   accident_year                         189815 non-null  int64
2   accident_reference                   189815 non-null  object
3   vehicle_reference                    189815 non-null  int64
4   vehicle_type                         189815 non-null  int64
5   towing_and_articulation              189815 non-null  int64
6   vehicle_manoeuvre                   189815 non-null  int64
7   vehicle_direction_from               189815 non-null  int64
8   vehicle_direction_to                 189815 non-null  int64
9   vehicle_location_restricted_lane     189815 non-null  int64
10  junction_location                   189815 non-null  int64
11  skidding_and_overturning             189815 non-null  int64
12  hit_object_in_carriageway            189815 non-null  int64
13  vehicle_leaving_carriageway          189815 non-null  int64
14  hit_object_off_carriageway           189815 non-null  int64
15  first_point_of_impact                 189815 non-null  int64
16  vehicle_left_hand_drive              189815 non-null  int64
17  journey_purpose_of_driver              189815 non-null  int64
18  sex_of_driver                        189815 non-null  int64
19  age_of_driver                        189815 non-null  int64
20  age_band_of_driver                  189815 non-null  int64
21  engine_capacity_cc                  189815 non-null  int64
22  propulsion_code                     189815 non-null  int64
23  age_of_vehicle                      189815 non-null  int64
24  generic_make_model                  189815 non-null  object
25  driver_imd_decile                   189815 non-null  int64
26  driver_home_area_type               189815 non-null  int64
27  lsoa_of_driver                      189815 non-null  object
28  escooter_flag                       189815 non-null  int64
29  dir_from_e                          19527 non-null   float64
30  dir_from_n                          19527 non-null   float64
31  dir_to_e                            19363 non-null   float64
32  dir_to_n                            19363 non-null   float64
33  driver_distance_banding              189815 non-null  int64
dtypes: float64(4), int64(26), object(4)
memory usage: 49.2+ MB
```

```
[11]: #Verify the data is loaded in correctly
casualty_data.head()
```

```
[11]:  accident_index  accident_year accident_reference  vehicle_reference \
0  2023010419171      2023      10419171      1
1  2023010419183      2023      10419183      2
2  2023010419183      2023      10419183      3
3  2023010419189      2023      10419189      1
4  2023010419191      2023      10419191      2

  casualty_reference  casualty_class  sex_of_casualty  age_of_casualty \
0          1          3          2          20
1          1          1          1          25
2          2          2          2          38
3          1          1          1          50
4          1          1          1          34

  age_band_of_casualty  casualty_severity  ...  pedestrian_movement \
0          4          3  ...          1
1          5          3  ...          0
2          7          3  ...          0
3          8          3  ...          0
4          6          3  ...          0

  car_passenger  bus_or_coach_passenger  pedestrian_road_maintenance_worker \
0          0          0          0
1          0          0          0
2          2          0          0
3          0          0          0
4          0          0          0

  casualty_type  casualty_home_area_type  casualty_imd_decile \
0          0          1          10
1          9          1          3
2          9          -1          -1
3          9          1          5
4          1          1          2

  lsoa_of_casualty  enhanced_casualty_severity  casualty_distance_banding
0  E01030370          -1          3
1  E01001546          -1          1
2          -1          -1          -1
3  E01002443          -1          1
4  E01004679          -1          2
```

```
[5 rows x 21 columns]
```

```
[12]: #Verify the total number of samples and features of the data
casualty_data.shape
```

```
[12]: (132977, 21)
```

```
[13]: #Check the information of the dataset
casualty_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 132977 entries, 0 to 132976
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   accident_index                       132977 non-null  object
1   accident_year                        132977 non-null  int64
2   accident_reference                   132977 non-null  object
3   vehicle_reference                   132977 non-null  int64
4   casualty_reference                   132977 non-null  int64
5   casualty_class                       132977 non-null  int64
6   sex_of_casualty                     132977 non-null  int64
7   age_of_casualty                     132977 non-null  int64
8   age_band_of_casualty                 132977 non-null  int64
9   casualty_severity                   132977 non-null  int64
10  pedestrian_location                  132977 non-null  int64
11  pedestrian_movement                  132977 non-null  int64
12  car_passenger                       132977 non-null  int64
13  bus_or_coach_passenger               132977 non-null  int64
14  pedestrian_road_maintenance_worker  132977 non-null  int64
15  casualty_type                       132977 non-null  int64
16  casualty_home_area_type              132977 non-null  int64
17  casualty_imd_decile                  132977 non-null  int64
18  lsoa_of_casualty                    132977 non-null  object
19  enhanced_casualty_severity           132977 non-null  int64
20  casualty_distance_banding            132977 non-null  int64
dtypes: int64(18), object(3)
memory usage: 21.3+ MB
```

```
[14]: #Verify the data is loaded in correctly
collision_data.head()
```

```
[14]:  accident_index  accident_year  accident_reference  location_easting_osgr  \
0   2023010419171          2023          10419171          525060.0
1   2023010419183          2023          10419183          535463.0
2   2023010419189          2023          10419189          508702.0
3   2023010419191          2023          10419191          520341.0
4   2023010419192          2023          10419192          527255.0

    location_northing_osgr  longitude  latitude  police_force  \
```

0	170416.0	-0.202878	51.418974	1
1	198745.0	-0.042464	51.671155	1
2	177696.0	-0.435789	51.487777	1
3	190175.0	-0.263972	51.597575	1
4	176963.0	-0.168976	51.477324	1

	accident_severity	number_of_vehicles	...	light_conditions	\
0	3		1 ...	4	
1	3		3 ...	4	
2	3		2 ...	4	
3	3		2 ...	4	
4	3		2 ...	4	

	weather_conditions	road_surface_conditions	special_conditions_at_site	\
0	8		2	0
1	1		1	0
2	1		1	0
3	9		1	0
4	1		1	0

	carriageway_hazards	urban_or_rural_area	\
0	0	1	
1	0	1	
2	0	1	
3	0	1	
4	0	1	

	did_police_officer_attend_scene_of_accident	trunk_road_flag	\
0		1	2
1		1	2
2		1	2
3		1	2
4		1	2

	lsoa_of_accident_location	enhanced_severity_collision
0	E01003383	-1
1	E01001547	-1
2	E01002448	-1
3	E01000129	-1
4	E01004583	-1

[5 rows x 37 columns]

```
[15]: #Verify the total number of samples and features of the data
collision_data.shape
```

```
[15]: (104258, 37)
```

```
[16]: #Check the information of the dataset
collision_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 104258 entries, 0 to 104257
Data columns (total 37 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   accident_index                           104258 non-null  object
1   accident_year                            104258 non-null  int64
2   accident_reference                       104258 non-null  object
3   location_easting_osgr                   104246 non-null  float64
4   location_northing_osgr                  104246 non-null  float64
5   longitude                                104246 non-null  float64
6   latitude                                 104246 non-null  float64
7   police_force                            104258 non-null  int64
8   accident_severity                       104258 non-null  int64
9   number_of_vehicles                      104258 non-null  int64
10  number_of_casualties                    104258 non-null  int64
11  date                                    104258 non-null  object
12  day_of_week                             104258 non-null  int64
13  time                                    104258 non-null  object
14  local_authority_district                 104258 non-null  int64
15  local_authority_ons_district             104258 non-null  object
16  local_authority_highway                  104258 non-null  object
17  first_road_class                         104258 non-null  int64
18  first_road_number                       104258 non-null  int64
19  road_type                               104258 non-null  int64
20  speed_limit                             104258 non-null  int64
21  junction_detail                         104258 non-null  int64
22  junction_control                        104258 non-null  int64
23  second_road_class                       104258 non-null  int64
24  second_road_number                      104258 non-null  int64
25  pedestrian_crossing_human_control        104258 non-null  int64
26  pedestrian_crossing_physical_facilities  104258 non-null  int64
27  light_conditions                        104258 non-null  int64
28  weather_conditions                      104258 non-null  int64
29  road_surface_conditions                  104258 non-null  int64
30  special_conditions_at_site               104258 non-null  int64
31  carriageway_hazards                     104258 non-null  int64
32  urban_or_rural_area                     104258 non-null  int64
33  did_police_officer_attend_scene_of_accident 104258 non-null  int64
34  trunk_road_flag                         104258 non-null  int64
35  lsoa_of_accident_location               104258 non-null  object
36  enhanced_severity_collision              104258 non-null  int64
dtypes: float64(4), int64(26), object(7)
memory usage: 29.4+ MB
```


3.2 Merging the Three Datasets

```
[18]: #The three datasets were merged using accident_index, with the unique value for
      ↪ each accident acting as the key.
      #Merge the Casualty and the Vehicle dataset first
      accident_1 = pd.merge(casualty_data, vehicle_data, on = 'accident_index',
      ↪ how='inner')

      #Merge the result with the Collision dataset:
      accident = pd.merge(accident_1, collision_data, on = 'accident_index',
      ↪ how='inner')
      accident.head()
```

```
[18]:  accident_index  accident_year_x  accident_reference_x  vehicle_reference_x  \
0    2023010419171          2023          10419171          1
1    2023010419183          2023          10419183          2
2    2023010419183          2023          10419183          2
3    2023010419183          2023          10419183          2
4    2023010419183          2023          10419183          3

      casualty_reference  casualty_class  sex_of_casualty  age_of_casualty  \
0              1              3              2          20
1              1              1              1          25
2              1              1              1          25
3              1              1              1          25
4              2              2              2          38

      age_band_of_casualty  casualty_severity  ...  light_conditions  \
0              4              3  ...              4
1              5              3  ...              4
2              5              3  ...              4
3              5              3  ...              4
4              7              3  ...              4

      weather_conditions  road_surface_conditions  special_conditions_at_site  \
0              8              2              0
1              1              1              0
2              1              1              0
3              1              1              0
4              1              1              0

      carriageway_hazards  urban_or_rural_area  \
0              0              1
1              0              1
2              0              1
3              0              1
4              0              1
```

	did_police_officer_attend_scene_of_accident	trunk_road_flag \
0	1	2
1	1	2
2	1	2
3	1	2
4	1	2

	lsoa_of_accident_location	enhanced_severity_collision
0	E01003383	-1
1	E01001547	-1
2	E01001547	-1
3	E01001547	-1
4	E01001547	-1

[5 rows x 90 columns]

```
[19]: #Check the information of the new merged dataset
accident.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 161607 entries, 0 to 161606
Data columns (total 90 columns):
```

#	Column	Non-Null Count	Dtype
0	accident_index	161607 non-null	object
1	accident_year_x	161607 non-null	int64
2	accident_reference_x	161607 non-null	object
3	vehicle_reference_x	161607 non-null	int64
4	casualty_reference	161607 non-null	int64
5	casualty_class	161607 non-null	int64
6	sex_of_casualty	161607 non-null	int64
7	age_of_casualty	161607 non-null	int64
8	age_band_of_casualty	161607 non-null	int64
9	casualty_severity	161607 non-null	int64
10	pedestrian_location	161607 non-null	int64
11	pedestrian_movement	161607 non-null	int64
12	car_passenger	161607 non-null	int64
13	bus_or_coach_passenger	161607 non-null	int64
14	pedestrian_road_maintenance_worker	161607 non-null	int64
15	casualty_type	161607 non-null	int64
16	casualty_home_area_type	161607 non-null	int64
17	casualty_imd_decile	161607 non-null	int64
18	lsoa_of_casualty	161607 non-null	object
19	enhanced_casualty_severity	161607 non-null	int64
20	casualty_distance_banding	161607 non-null	int64
21	accident_year_y	161607 non-null	int64
22	accident_reference_y	161607 non-null	object

23	vehicle_reference_y	161607	non-null	int64
24	vehicle_type	161607	non-null	int64
25	towing_and_articulation	161607	non-null	int64
26	vehicle_manoeuvre	161607	non-null	int64
27	vehicle_direction_from	161607	non-null	int64
28	vehicle_direction_to	161607	non-null	int64
29	vehicle_location_restricted_lane	161607	non-null	int64
30	junction_location	161607	non-null	int64
31	skidding_and_overturning	161607	non-null	int64
32	hit_object_in_carriageway	161607	non-null	int64
33	vehicle_leaving_carriageway	161607	non-null	int64
34	hit_object_off_carriageway	161607	non-null	int64
35	first_point_of_impact	161607	non-null	int64
36	vehicle_left_hand_drive	161607	non-null	int64
37	journey_purpose_of_driver	161607	non-null	int64
38	sex_of_driver	161607	non-null	int64
39	age_of_driver	161607	non-null	int64
40	age_band_of_driver	161607	non-null	int64
41	engine_capacity_cc	161607	non-null	int64
42	propulsion_code	161607	non-null	int64
43	age_of_vehicle	161607	non-null	int64
44	generic_make_model	161607	non-null	object
45	driver_imd_decile	161607	non-null	int64
46	driver_home_area_type	161607	non-null	int64
47	lsoa_of_driver	161607	non-null	object
48	escooter_flag	161607	non-null	int64
49	dir_from_e	12355	non-null	float64
50	dir_from_n	12355	non-null	float64
51	dir_to_e	12236	non-null	float64
52	dir_to_n	12236	non-null	float64
53	driver_distance_banding	161607	non-null	int64
54	accident_year	161607	non-null	int64
55	accident_reference	161607	non-null	object
56	location_easting_osgr	161591	non-null	float64
57	location_northing_osgr	161591	non-null	float64
58	longitude	161591	non-null	float64
59	latitude	161591	non-null	float64
60	police_force	161607	non-null	int64
61	accident_severity	161607	non-null	int64
62	number_of_vehicles	161607	non-null	int64
63	number_of_casualties	161607	non-null	int64
64	date	161607	non-null	object
65	day_of_week	161607	non-null	int64
66	time	161607	non-null	object
67	local_authority_district	161607	non-null	int64
68	local_authority_ons_district	161607	non-null	object
69	local_authority_highway	161607	non-null	object
70	first_road_class	161607	non-null	int64

```

71 first_road_number          161607 non-null int64
72 road_type                  161607 non-null int64
73 speed_limit                161607 non-null int64
74 junction_detail            161607 non-null int64
75 junction_control            161607 non-null int64
76 second_road_class           161607 non-null int64
77 second_road_number          161607 non-null int64
78 pedestrian_crossing_human_control 161607 non-null int64
79 pedestrian_crossing_physical_facilities 161607 non-null int64
80 light_conditions            161607 non-null int64
81 weather_conditions          161607 non-null int64
82 road_surface_conditions      161607 non-null int64
83 special_conditions_at_site    161607 non-null int64
84 carriageway_hazards          161607 non-null int64
85 urban_or_rural_area          161607 non-null int64
86 did_police_officer_attend_scene_of_accident 161607 non-null int64
87 trunk_road_flag             161607 non-null int64
88 lsoa_of_accident_location     161607 non-null object
89 enhanced_severity_collision    161607 non-null int64
dtypes: float64(8), int64(70), object(12)
memory usage: 111.0+ MB

```

3.3 Dropping Columns

```
[21]: #Examine the available columns that can be dropped
accident.columns
```

```
[21]: Index(['accident_index', 'accident_year_x', 'accident_reference_x',
'vehicle_reference_x', 'casualty_reference', 'casualty_class',
'sex_of_casualty', 'age_of_casualty', 'age_band_of_casualty',
'casualty_severity', 'pedestrian_location', 'pedestrian_movement',
'car_passenger', 'bus_or_coach_passenger',
'pedestrian_road_maintenance_worker', 'casualty_type',
'casualty_home_area_type', 'casualty_imd_decile', 'lsoa_of_casualty',
'enhanced_casualty_severity', 'casualty_distance_banding',
'accident_year_y', 'accident_reference_y', 'vehicle_reference_y',
'vehicle_type', 'towing_and_articulation', 'vehicle_manoeuvre',
'vehicle_direction_from', 'vehicle_direction_to',
'vehicle_location_restricted_lane', 'junction_location',
'skidding_and_overturning', 'hit_object_in_carriageway',
'vehicle_leaving_carriageway', 'hit_object_off_carriageway',
'first_point_of_impact', 'vehicle_left_hand_drive',
'journey_purpose_of_driver', 'sex_of_driver', 'age_of_driver',
'age_band_of_driver', 'engine_capacity_cc', 'propulsion_code',
'age_of_vehicle', 'generic_make_model', 'driver_imd_decile',
'driver_home_area_type', 'lsoa_of_driver', 'escooter_flag',
'dir_from_e', 'dir_from_n', 'dir_to_e', 'dir_to_n',
```

```

'driver_distance_banding', 'accident_year', 'accident_reference',
'location_easting_osgr', 'location_northing_osgr', 'longitude',
'latitude', 'police_force', 'accident_severity', 'number_of_vehicles',
'number_of_casualties', 'date', 'day_of_week', 'time',
'local_authority_district', 'local_authority_ons_district',
'local_authority_highway', 'first_road_class', 'first_road_number',
'road_type', 'speed_limit', 'junction_detail', 'junction_control',
'second_road_class', 'second_road_number',
'pedestrian_crossing_human_control',
'pedestrian_crossing_physical_facilities', 'light_conditions',
'weather_conditions', 'road_surface_conditions',
'special_conditions_at_site', 'carriageway_hazards',
'urban_or_rural_area', 'did_police_officer_attend_scene_of_accident',
'trunk_road_flag', 'lsoa_of_accident_location',
'enhanced_severity_collision'],
dtype='object')

```

```

[22]: #Columns were dropped based on whether the target variable accident severity
      ↪ were influenced by them or not.
      #The columns below were the ones that were picked.
main_accident = accident.loc[:, ['accident_index', 'time', 'number_of_vehicles',
      ↪ 'number_of_casualties', 'date',
      ↪ 'local_authority_ons_district',
      ↪ 'road_type', 'speed_limit', 'junction_detail',
      ↪ 'junction_control',
      ↪ 'light_conditions', 'weather_conditions',
      ↪ 'road_surface_conditions',
      ↪
      ↪ 'urban_or_rural_area', 'vehicle_location_restricted_lane',
      ↪ 'junction_location',
      ↪ 'skidding_and_overturning',
      ↪ 'hit_object_in_carriageway',
      ↪ 'vehicle_leaving_carriageway',
      ↪ 'hit_object_off_carriageway', 'first_point_of_impact',
      ↪ 'journey_purpose_of_driver',
      ↪ 'sex_of_driver', 'age_of_driver',
      ↪ 'vehicle_type', 'vehicle_manoeuvre',
      ↪ 'casualty_type',
      ↪ 'accident_severity']]

```

```

[23]: #Verify that the columns were kept
main_accident.columns

```

```

[23]: Index(['accident_index', 'time', 'number_of_vehicles', 'number_of_casualties',
      'date', 'local_authority_ons_district', 'road_type', 'speed_limit',
      'junction_detail', 'junction_control', 'light_conditions',

```

```

'weather_conditions', 'road_surface_conditions', 'urban_or_rural_area',
'vehicle_location_restricted_lane', 'junction_location',
'skidding_and_overturning', 'hit_object_in_carriageway',
'vehicle_leaving_carriageway', 'hit_object_off_carriageway',
'first_point_of_impact', 'journey_purpose_of_driver', 'sex_of_driver',
'age_of_driver', 'vehicle_type', 'vehicle_manoeuvre', 'casualty_type',
'accident_severity'],
dtype='object')

```

3.4 Variable Description

Variables are categorised into 3 main types:

- **Numeric:** Variables containing numeric values.
- **Categorical:** Variables containing text data / each unique value indicates a category.
- **Date/Time:** Variables containing date or time values, like accident date or timestamp.

No.	Input data	Definition	Category
1	accident_index	Unique reference for the accident	Categorical
2	number_of_vehicles	Number of vehicles involved in the accident	Numeric
3	number_of_casualties	Number of casualties in the accident	Numeric
4	date	Date of the accident (in Date/Time format)	Date/Time
5	local_authority_ons_district	Local authority districts where the accident occurred using codes from Office for National Statistics (ONS)	Categorical
6	road_type	Type of road where the accident occurred	Categorical
7	speed_limit	Speed limit on the road at the accident location	Numeric
8	junction_detail	Details about the junction at the accident location	Categorical
9	junction_control	Type of control at the junction	Categorical
10	light_conditions	Light conditions at the time of the accident	Categorical
11	weather_conditions	Weather conditions during the accident	Categorical
12	road_surface_conditions	Road surface conditions at the accident location	Categorical
13	urban_or_rural_area	Whether the accident occurred in an urban or rural area	Categorical
14	vehicle_location_restricted_lane	Whether the vehicle was in a restricted lane at the time of the accident	Categorical
15	junction_location	Location of the junction where the accident occurred	Categorical
16	skidding_and_overturning	Whether the vehicle skidded or overturned during the accident	Categorical
17	hit_object_in_carriageway	Whether the vehicle hit an object in the carriageway	Categorical

No.	Input data	Definition	Category
18	vehicle_leaving_carriageway	Whether the vehicle left the carriageway during the accident	Categorical
19	hit_object_off_carriageway	Whether the vehicle hit an object off the carriageway	Categorical
20	first_point_of_impact	First point of impact in the accident	Categorical
21	journey_purpose_of_driver	Purpose of the driver's journey at the time of the accident	Categorical
22	sex_of_driver	Gender of the driver involved in the accident	Categorical
23	age_of_driver	Age of the driver at the time of the accident	Numeric
24	vehicle_type	Type of vehicle involved in the accident	Categorical
25	vehicle_manoeuvre	Manoeuvre the vehicle was performing at the time of the accident	Categorical
26	casualty_type	Type of casualty	Categorical
27	accident_severity	Severity of the casualty	Categorical

```
[25]: #Filter the main_accident dataset to only include accidents occurred in London

#These are rows that the local_authority_ons_district column starts with the
↳prefix 'E09'
main_accident_filtered =
↳main_accident[main_accident['local_authority_ons_district'].str.
↳startswith('E09', na=False)]
```

```
[26]: #Verify the samples and features of the dataset after the filtering
main_accident_filtered.shape
```

```
[26]: (48149, 28)
```

```
[27]: #Drop duplicates
main_accident_filtered.drop_duplicates(keep = 'first', inplace = True)
```

```
[28]: #Verify the samples and features of the dataset after dropping duplicates
main_accident_filtered.shape
```

```
[28]: (42129, 28)
```

```
[29]: #Check example samples
main_accident_filtered.head()
```

```
[29]:  accident_index  time  number_of_vehicles  number_of_casualties  date \
0  2023010419171  01:24                1                1  01/01/2023
1  2023010419183  02:25                3                2  01/01/2023
2  2023010419183  02:25                3                2  01/01/2023
3  2023010419183  02:25                3                2  01/01/2023
7  2023010419189  03:50                2                1  01/01/2023
```

	local_authority_ons_district	road_type	speed_limit	junction_detail	\
0	E09000024	2	20	9	
1	E09000010	6	30	3	
2	E09000010	6	30	3	
3	E09000010	6	30	3	
7	E09000017	1	30	1	

	junction_control	...	vehicle_leaving_carriageway	\
0	4	...	0	
1	4	...	0	
2	4	...	0	
3	4	...	0	
7	4	...	0	

	hit_object_off_carriageway	first_point_of_impact	\
0	0	1	
1	0	1	
2	0	4	
3	0	1	
7	0	4	

	journey_purpose_of_driver	sex_of_driver	age_of_driver	vehicle_type	\
0	1	1	61	11	
1	1	1	54	11	
2	6	1	25	9	
3	6	1	42	9	
7	6	1	50	9	

	vehicle_manoeuvre	casualty_type	accident_severity
0	4	0	3
1	18	9	3
2	9	9	3
3	8	9	3
7	18	9	3

[5 rows x 28 columns]

```
[30]: #Count the number of accidents in each London district
main_accident_filtered['local_authority_ons_district'].value_counts()
```

```
[30]: local_authority_ons_district
E09000033    2223
E09000008    2080
E09000032    1963
E09000022    1908
E09000028    1889
```



```

E09000010    1881
E09000030    1868
E09000014    1797
E09000005    1663
E09000009    1579
E09000003    1565
E09000025    1475
E09000023    1464
E09000012    1397
E09000007    1290
E09000018    1251
E09000017    1189
E09000020    1178
E09000026    1164
E09000011    1132
E09000006    1131
E09000013    1115
E09000019    1065
E09000031     941
E09000016     936
E09000002     837
E09000024     823
E09000004     743
E09000015     723
E09000029     683
E09000027     668
E09000021     508
Name: count, dtype: int64

```

```

[31]: #Top 5 boroughs with the highest number of accidents are:
boroughs = ['E09000033', 'E09000008', 'E09000032', 'E09000022', 'E09000028']

#Filter the 'main_accident_filtered' dataframe to include only rows for these
↳ top 5 boroughs
top_5_boroughs = main_accident_filtered.
↳ loc[main_accident_filtered['local_authority_ons_district'].isin(boroughs)]

#Display the summary information about the filtered dataframe
top_5_boroughs.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Index: 10063 entries, 11 to 48187
Data columns (total 28 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   accident_index        10063 non-null  object
 1   time                  10063 non-null  object
 2   number_of_vehicles     10063 non-null  int64

```

3	number_of_casualties	10063	non-null	int64
4	date	10063	non-null	object
5	local_authority_ons_district	10063	non-null	object
6	road_type	10063	non-null	int64
7	speed_limit	10063	non-null	int64
8	junction_detail	10063	non-null	int64
9	junction_control	10063	non-null	int64
10	light_conditions	10063	non-null	int64
11	weather_conditions	10063	non-null	int64
12	road_surface_conditions	10063	non-null	int64
13	urban_or_rural_area	10063	non-null	int64
14	vehicle_location_restricted_lane	10063	non-null	int64
15	junction_location	10063	non-null	int64
16	skidding_and_overturning	10063	non-null	int64
17	hit_object_in_carriageway	10063	non-null	int64
18	vehicle_leaving_carriageway	10063	non-null	int64
19	hit_object_off_carriageway	10063	non-null	int64
20	first_point_of_impact	10063	non-null	int64
21	journey_purpose_of_driver	10063	non-null	int64
22	sex_of_driver	10063	non-null	int64
23	age_of_driver	10063	non-null	int64
24	vehicle_type	10063	non-null	int64
25	vehicle_manoeuvre	10063	non-null	int64
26	casualty_type	10063	non-null	int64
27	accident_severity	10063	non-null	int64

dtypes: int64(24), object(4)
memory usage: 2.2+ MB

```
[32]: #Convert the object type of the date and time columns into a datetime type.
top_5_boroughs["date"] = pd.to_datetime(top_5_boroughs["date"], format="%d/%m/%Y")
top_5_boroughs["time"] = pd.to_datetime(top_5_boroughs["time"], format="%H:%M")

#Check to see if the changes were made
print(top_5_boroughs.dtypes)
```

accident_index	object
time	datetime64[ns]
number_of_vehicles	int64
number_of_casualties	int64
date	datetime64[ns]
local_authority_ons_district	object
road_type	int64
speed_limit	int64
junction_detail	int64
junction_control	int64
light_conditions	int64
weather_conditions	int64

```

road_surface_conditions          int64
urban_or_rural_area             int64
vehicle_location_restricted_lane int64
junction_location              int64
skidding_and_overturning        int64
hit_object_in_carriageway       int64
vehicle_leaving_carriageway     int64
hit_object_off_carriageway      int64
first_point_of_impact           int64
journey_purpose_of_driver         int64
sex_of_driver                   int64
age_of_driver                   int64
vehicle_type                    int64
vehicle_manoeuvre              int64
casualty_type                   int64
accident_severity               int64
dtype: object

```

```

[33]: #Split the original dataframe into two different dataframes
#The X dataframe has all the variables except the target variable
X = top_5_boroughs.loc[:, [col for col in top_5_boroughs.columns if col !=
    ↪ 'accident_severity']]

#The y dataframe contains the target variable
y = top_5_boroughs.loc[:, 'accident_severity']

```

```

[34]: #Check the shape of the two dataframes
print(X.shape)
print(y.shape)

```

```

(10063, 27)
(10063,)

```

```

[35]: #Check for any missing values in both the X and y dataframes. -1 is a stand-in
    ↪ value to represent missing values.

search_value = -1

#Search for any -1 value
X_columns_with_value = X.columns[(X == search_value).any(axis=0)]
has_negative_one = (y == -1).any().any()

#Print the values
print(f"The following columns in the X dataframe have -1 values:
    ↪ \n{X_columns_with_value}\n")
print(f"There are -1 values in the y dataframe: {has_negative_one}")

```

The following columns in the X dataframe have -1 values:

```
Index(['junction_control', 'age_of_driver'], dtype='object')
```

There are -1 values in the y dataframe: False

```
[36]: #Change all the -1 values into nan values to show that they represent missing
      ↪ values
      X = X.replace(-1, np.nan)

      #Check if the changes were implemented
      X.isnull().sum()
```

```
[36]: accident_index          0
      time                    0
      number_of_vehicles      0
      number_of_casualties    0
      date                    0
      local_authority_ons_district  0
      road_type               0
      speed_limit             0
      junction_detail         0
      junction_control        1718
      light_conditions        0
      weather_conditions      0
      road_surface_conditions  0
      urban_or_rural_area     0
      vehicle_location_restricted_lane  0
      junction_location       0
      skidding_and_overturning  0
      hit_object_in_carriageway  0
      vehicle_leaving_carriageway  0
      hit_object_off_carriageway  0
      first_point_of_impact    0
      journey_purpose_of_driver  0
      sex_of_driver           0
      age_of_driver           2492
      vehicle_type            0
      vehicle_manoeuvre       0
      casualty_type           0
      dtype: int64
```

4 Data Splitting

4.1 Class Distribution of Target Variable


Split the X and Y datasets into training and test sets. The y_train and y_test sets only have the target variable.

```
[39]: #Calculate percentage proportions
percentage_proportions = top_5_boroughs['accident_severity'].
    ↪value_counts(normalize=True) * 100
#Format each value in the Series using lamda
formatted_percentages = percentage_proportions.apply(lambda x: f'{x:.2f}%')
print(formatted_percentages)
```

```
accident_severity
3    84.11%
2    15.62%
1     0.27%
Name: proportion, dtype: object
```

```
[40]: #Highlight the number of counts for each level of severity in the target_
    ↪variable
top_5_boroughs["accident_severity"].value_counts().plot(kind="bar", rot=0)
```

```
[40]: <Axes: xlabel='accident_severity'>
```



As a result of the uneven distribution of the classes within the target variable using stratifiedshufflesplit, this is because using random sampling may introduce a bias into the proportion.

```
[42]: #Split the data from the two dataframes into smaller subsets; there should be a
      ↪training set and a testing set for each of the two dataframes.
      #Reset indices in case the two dataframes are not aligned.
      X = X.reset_index(drop=True)
      y = y.reset_index(drop=True)

      #The test set should be 20% of the total dataset while the training set should
      ↪be 80%.
      stratified_splitter = StratifiedShuffleSplit(n_splits=1, test_size=0.2,
      ↪random_state=7)
      train_index, test_index = next(stratified_splitter.split(X, y))

      #Apply indices to both X and y
      X_train, X_test = X.iloc[train_index], X.iloc[test_index]
      y_train, y_test = y.iloc[train_index], y.iloc[test_index]

      #Show the shape of the four dataframes
      print(X_train.shape)
      print(X_test.shape)
      print(y_train.shape)
      print(y_test.shape)

      (8050, 27)
      (2013, 27)
      (8050,)
      (2013,)
```

```
[43]: #Create a function that measures the proportions of an accident's severity
      def accident_severity_proportions(data):
          target = data["accident_severity"] if isinstance(data, pd.DataFrame) else
          ↪data
          return target.value_counts() / len(data)

      #Create a random split
      rand_train_set, rand_test_set = train_test_split(top_5_boroughs, test_size=0.2,
      ↪random_state=7)

      #Create a temporary dataframe for easy visualization
      df_tmp = pd.DataFrame({
          "Overall": accident_severity_proportions(top_5_boroughs),
          "Random test set": accident_severity_proportions(rand_test_set),
          "Stratified test set": accident_severity_proportions(y_test),
      }).sort_index()

      #Add two columns for the percent of the difference to the overall proportion
```

```

df_tmp["Rand. %error"] = 100 * df_tmp["Random test set"] / df_tmp["Overall"] - 100
df_tmp["Strat. %error"] = 100 * df_tmp["Stratified test set"] / df_tmp["Overall"] - 100

#Visualize table
df_tmp

```

```

[43]:

```

	Overall	Random test set	Stratified test set	\
accident_severity				
1	0.002683	0.001987	0.002484	
2	0.156216	0.156483	0.156483	
3	0.841101	0.841530	0.841033	

	Rand. %error	Strat. %error
accident_severity		
1	-25.940645	-7.425806
2	0.170931	0.170931
3	0.051004	-0.008058

5 Exploratory Data Analysis

Exploratory data analysis on the training set includes descriptive statistics and visualisations for numerical and categorical predictors to identify distributions, outliers, and value counts.

5.1 Univariate Analysis

5.1.1 Target Variable

```

[48]: #Check the target variables training dataset
y_train.head()

```

```

[48]: 2142    3
      1930    3
      8757    3
      9327    3
      7125    2
      Name: accident_severity, dtype: int64

```

```

[49]: #Descriptive Statistics of the Target Variable

#Frequency counts
frequency_counts = pd.DataFrame(y_train.value_counts().sort_index())

#Mode
mode_value = y_train.mode()[0]

```

```

#Median
median_value = y_train.median()

#Percentiles (25th, 50th, 75th)
percentiles = pd.DataFrame(y_train.quantile([0.25, 0.5, 0.75]))

#Summary
print("Frequency Counts:\n", frequency_counts)
print("\nMode:", mode_value)
print("Median:", median_value)
print("\nPercentiles (25th, 50th, 75th):\n", percentiles)

```

Frequency Counts:

	count
accident_severity	
1	22
2	1257
3	6771

Mode: 3

Median: 3.0

Percentiles (25th, 50th, 75th):

	accident_severity
0.25	3.0
0.50	3.0
0.75	3.0

The mode of the accident severity is 3, which means that most of the accidents had a slight severity, backed up by the frequencies of the accident severity and the percentile table indicates that 75% of the accidents have a severity rating of slight.

```

[51]: #Set the size of the figure
plt.figure(figsize=(10, 6))

#Visualise the distribution of the target variables
ax = sns.countplot(x=y_train, order=y.value_counts().index, palette='Reds')

#Add counts on bars
for container in ax.containers:
    ax.bar_label(container, fontsize=10)

#Label the plot and create a legend
plt.title('Accident Severity Distribution', fontsize=14)
plt.legend(title='Accident Severity', labels=['3 - Slight', '2 - Serious', '1 - Fatal'],
           bbox_to_anchor=(1.05, 0.5), loc='center left')
plt.xlabel('Level of Severity')
plt.ylabel('Count')

```




```
plt.xticks(ha='center')
```

```
#Cleaner borders
```

```
sns.despine()
```

```
#Display the plot
```

```
plt.show()
```



The data suggests that while fatal accidents are the least frequent, there is a significant number of slight accidents, indicating a need for targeted safety interventions to reduce overall accident rates.

5.1.2 Predictor Variables

Numerical Variables

```
[55]: #There are four numerical predictor variables: Number of vehicles, number of  
      ↪ casualties, the speed limit, and the age of the drivers.
```

```
#Descriptive statistics
```

```
def q1(x):  
    return x.quantile(0.25)
```

```
def q3(x):  
    return x.quantile(0.75)
```

```
def iqr(x):  
    return q3(x) - q1(x)
```

```
def get_mode(x):  
    return x.mode().iloc[0]
```

```

#Calculate statistics with proper naming
stats = X_train[['number_of_vehicles', 'number_of_casualties', 'speed_limit',
↪ 'age_of_driver']].agg([
    'min',
    'max',
    q1,
    'median',
    q3,
    iqr,
    get_mode
]).rename(index={
    'q1': '25%',
    'median': '50%',
    'q3': '75%',
    'iqr': 'IQR',
    'get_mode': 'mode'
})

#Format and transpose results
stats_df = stats.transpose().reset_index().rename(columns={'index': 'variable'})
stats_df = stats_df[['variable', 'min', '25%', '50%', '75%', 'max', 'IQR',
↪ 'mode']]
stats_df = stats_df.set_index('variable')

print("Descriptive Statistics:")
display(stats_df)

```

Descriptive Statistics:

	min	25%	50%	75%	max	IQR	mode
variable							
number_of_vehicles	1.0	2.0	2.0	2.0	17.0	0.0	2.0
number_of_casualties	1.0	1.0	1.0	1.0	7.0	0.0	1.0
speed_limit	20.0	20.0	20.0	30.0	60.0	10.0	20.0
age_of_driver	8.0	28.0	36.0	48.0	93.0	20.0	33.0

number_of_vehicles

- As indicated in the table above and the plot below the highest number of vehicles involved in a single accident was 17.
- 75% of accidents consisted of 2 or less vehicles.

number_of_casualties

- The highest number of casualties in an accident was 7, however this is an outlier because the majority of accidents only had 1 casualty.
- The percentile ranges for casualties are 1, suggesting that most accidents involved only one individual facing harm.

speed_limit

- For the speed limit, the 25th and 50th percentiles are 20, while the 75th percentile is 30, suggesting that 75% of the speed limits are 30 or below.

age_of_driver

- Most drivers involved in the accidents were around 30.
- The maximum age of a driver in an accident was 93, while the minimum age was 5.

```
[60]: #Prepare and reshape the data into 3 columns using FacetGrid with the melt
      ↪function
      facet_data =
      ↪X_train[['number_of_vehicles', 'number_of_casualties', 'speed_limit']].melt()

      #Create the FacetGrid which plots facets by their 'variable' column
      g = sns.FacetGrid(facet_data, col="variable", col_wrap=3, height=6,
      ↪sharex=False, sharey=False)

      #Map the appropriate plot to each facet based on the variable type
      def plot_facet(variable, **kwargs):
          data = kwargs.pop('data')
          sns.countplot(x=data['value'], **kwargs)

      g.map_dataframe(plot_facet, 'value')


      #Remove gridlines from all plots using a for loop
      for ax in g.axes.flat:
          ax.grid(False)

      #Set the title for each facet dynamically to the column name
      g.set_titles("{col_name}")

      #Set labels for the x and y axes
      g.set_axis_labels("Value", "Frequency")

      #Main title for the entire plot and adjust its position
      g.fig.suptitle("Distribution of the Number of Vehicles, Number of Casualties,
      ↪and Speed Limit", fontsize=16, y=1.02)

      #Display the plots
      plt.show()
```



Most accidents involved two vehicles, at least one casualty, and occurred on 20 mph roads.

```
[62]: #Set the figure size
fig = plt.figure(figsize=(15, 5))

#Histogram
plt.subplot(1, 2, 1)

#Define plot object using histplot
hist = sns.histplot(X_train.loc[:, 'age_of_driver'], bins=25, kde=True)

#Setting graph title and labels
hist.set_title('Distribution of Age of the Drivers')
hist.set(xlabel='Ages', ylabel='Frequency')



#Boxplot
plt.subplot(1,2,2)

#Define plot object
box = sns.boxplot(X_train.loc[:, 'age_of_driver'])

#Setting graph title
box.set_title( 'Distribution of Age of the Drivers')
box.set(xlabel = 'Count', ylabel = 'Ages')

#Remove borders (spines)
sns.despine(fig=fig, right=True, top=True)

plt.show()
```



Most drivers were aged 25–35, with few outliers over 80.

Categorical Variables

```
[65]: #There are 20 categorical predictor variables.
#Converting all the numerical values into their actual values in the dataset
for column in X_train.columns:
    if column in guide_data['field name'].values:
        #Creating a dictionary to map code/format to label for each field_name
        mapping_dict = guide_data[guide_data['field name'] == column].
        ↪set_index('code/format')['label'].to_dict()
        #Replacing the values in top_5_boroughs based on the mapping dictionary
        if mapping_dict:
            X_train[column] = X_train[column].map(mapping_dict).
        ↪fillna(X_train[column])

for column in X_test.columns:
    if column in guide_data['field name'].values:
        #Creating a dictionary to map code/format to label for each field_name
        mapping_dict = guide_data[guide_data['field name'] == column].
        ↪set_index('code/format')['label'].to_dict()
        #Replacing the values in top_5_boroughs based on the mapping dictionary
        if mapping_dict:
            X_test[column] = X_test[column].map(mapping_dict).
        ↪fillna(X_test[column])

#Examine the newly converted data
X_train.head()
```

```
[65]:
```

	accident_index	time	number_of_vehicles	\
2142	2023010435768	1900-01-01 19:35:00	2.0	
1930	2023010434377	1900-01-01 15:09:00	1.0	
8757	2023010477755	1900-01-01 17:00:00	2.0	

9327	2023010481441	1900-01-01	17:15:00	2.0
7125	2023010467051	1900-01-01	18:40:00	2.0

	number_of_casualties	date	local_authority_ons_district	\
2142	1.0	2023-02-27	Lambeth	
1930	1.0	2023-03-24	Lambeth	
8757	1.0	2023-11-10	Lambeth	
9327	1.0	2023-12-01	Lambeth	
7125	1.0	2023-09-19	Wandsworth	

	road_type	speed_limit	junction_detail	\
2142	Unknown	30	Not at junction or within 20 metres	
1930	Unknown	20	Not at junction or within 20 metres	
8757	Single carriageway	30	T or staggered junction	
9327	Unknown	20	unknown (self reported)	
7125	Single carriageway	20	Not at junction or within 20 metres	

	junction_control	...	hit_object_in_carriageway	\
2142	NaN	...	unknown (self reported)	
1930	NaN	...	unknown (self reported)	
8757	Give way or uncontrolled	...	unknown (self reported)	
9327	unknown (self reported)	...	unknown (self reported)	
7125	NaN	...	0	

	vehicle_leaving_carriageway	hit_object_off_carriageway	\
2142	unknown (self reported)	unknown (self reported)	
1930	unknown (self reported)	unknown (self reported)	
8757	unknown (self reported)	unknown (self reported)	
9327	unknown (self reported)	unknown (self reported)	
7125	Did not leave carriageway	0	

	first_point_of_impact	journey_purpose_of_driver	sex_of_driver	\
2142	Front	Not known	Male	
1930	unknown (self reported)	Not known	Not known	
8757	unknown (self reported)	Journey as part of work	Male	
9327	unknown (self reported)	Not known	Female	
7125	Did not impact	Journey as part of work	Male	

	age_of_driver	vehicle_type	\
2142	26.0	Motorcycle 125cc and under	
1930	NaN	Car	
8757	25.0	Motorcycle 125cc and under	
9327	64.0	Car	
7125	60.0	Bus or coach (17 or more pass seats)	

	vehicle_manoeuvre	casualty_type
2142	unknown (self reported)	Motorcycle 125cc and under rider or passenger

1930	unknown (self reported)	Car occupant
8757	unknown (self reported)	Motorcycle 125cc and under rider or passenger
9327	unknown (self reported)	Pedestrian
7125	Waiting to go - held up	Bus or coach occupant (17 or more pass seats)

[5 rows x 27 columns]

```
[66]: #Create functions used for analyzing descriptive statistics and plot the
      ↪ categorical variable
```

```
#Function to describe the variable
def describe_column(data, column):
    # Get value counts and mode
    value_counts_result = data[column].value_counts().reset_index()
    value_counts_result.columns = ['Category', 'Frequency']
    mode_value = data[column].mode().iloc[0] # Get the first mode

    # Display the result
    print(f"Descriptive Statistics for {column}:")
    print(f"Mode: {mode_value}")
    display(value_counts_result)
    print("\n" + "-"*80) # Separator for readability

#Function to plot the variable
def plot_countplot(data, column):

    #Create figure
    plt.figure(figsize=(10, 6))

    #Create countplot
    ax = sns.countplot(data=data, x=column)

    #Apply customizations
    ax.set_xticklabels(ax.get_xticklabels(), rotation=45, ha='right')
    ax.grid(False)
    ax.set_title(f'Distribution of {column}', fontsize=14)
    ax.set_xlabel('Value', fontsize=12)
    ax.set_ylabel('Frequency', fontsize=12)

    #Remove plot borders
    for spine in ax.spines.values():
        spine.set_visible(False)

    #Adjust layout and show plot
    plt.tight_layout()
    plt.show()
```

local_authority_ons_district

```
[68]: #Use the previously created functions to describe and plot the categorical
      ↪variable
```

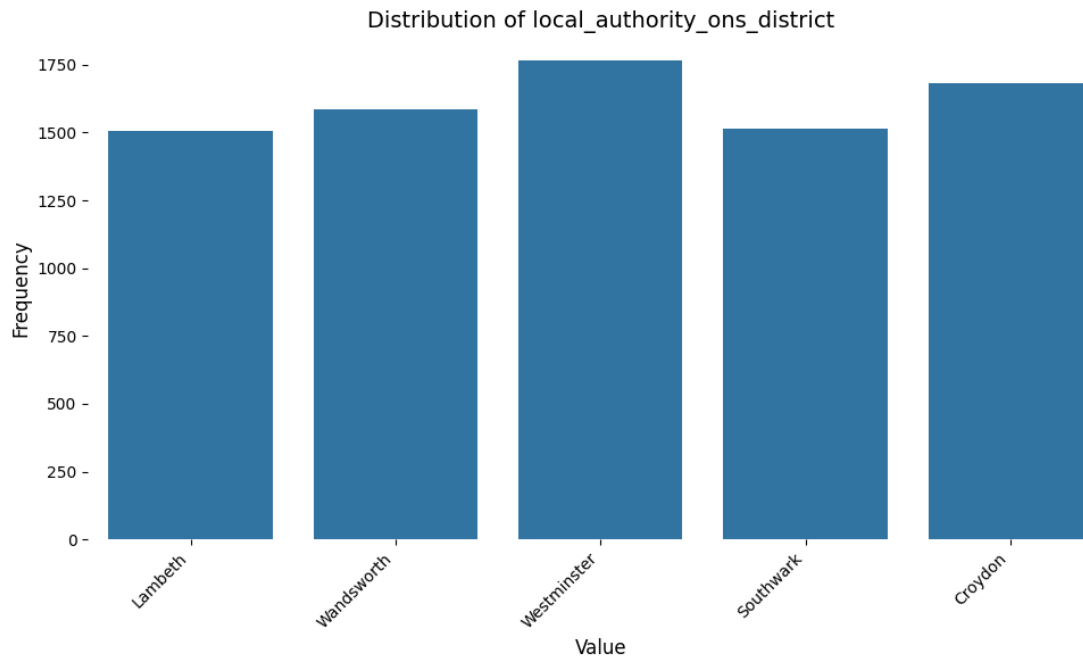
```
describe_column(X_train, 'local_authority_ons_district')
```

```
plot_countplot(X_train, 'local_authority_ons_district')
```

Descriptive Statistics for local_authority_ons_district:

Mode: Westminster

	Category	Frequency
0	Westminster	1764
1	Croydon	1681
2	Wandsworth	1587
3	Southwark	1513
4	Lambeth	1505



Accidents were most frequent in Westminster (1,764), with most boroughs reporting 1,505–1,764 incidents.

road_type

```
[71]: #Use the previously created functions to describe and plot the categorical
      ↪variable
```

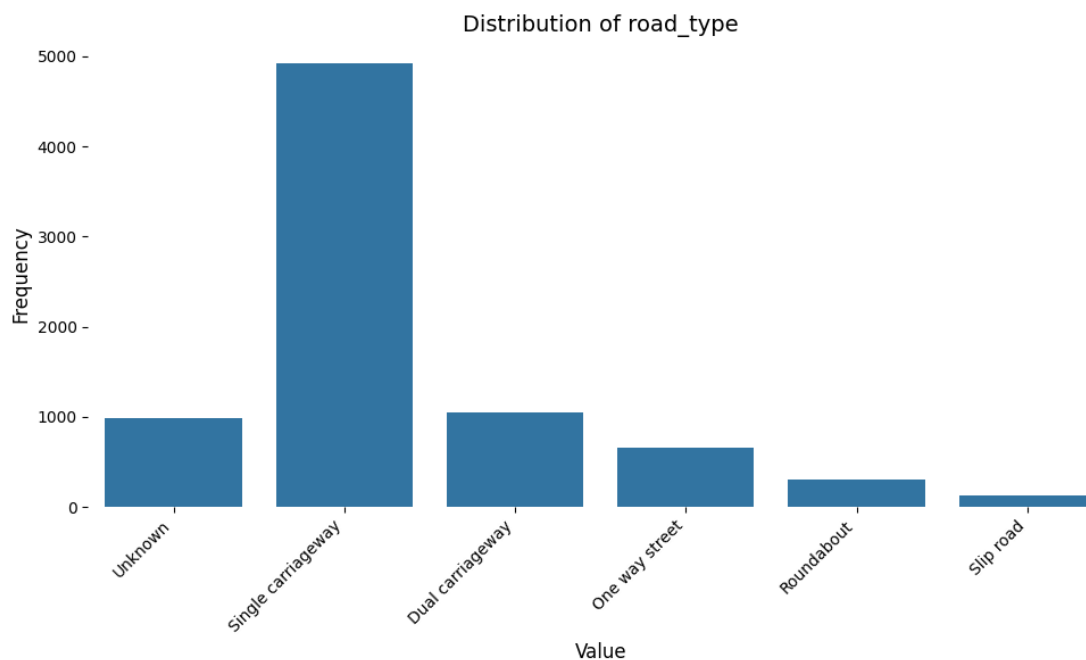
```
describe_column(X_train, 'road_type')
```

```
plot_countplot(X_train, 'road_type')
```


Descriptive Statistics for road_type:

Mode: Single carriageway

	Category	Frequency
0	Single carriageway	4923
1	Dual carriageway	1053
2	Unknown	984
3	One way street	657
4	Roundabout	303
5	Slip road	130



Single carriageways had the most accidents (4,923), suggesting higher risk due to traffic volume.

junction_detail

```
[74]: #Use the previously created functions to describe and plot the categorical_
      ↪variable
```

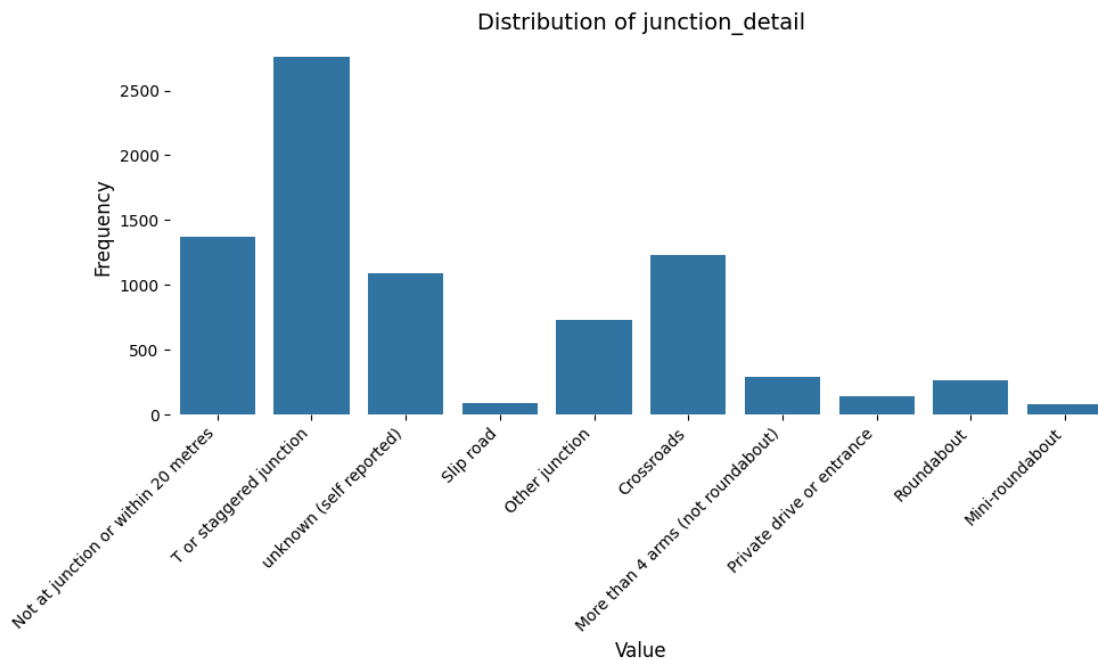
```
describe_column(X_train, 'junction_detail')
plot_countplot(X_train, 'junction_detail')
```

Descriptive Statistics for junction_detail:

Mode: T or staggered junction

	Category	Frequency
0	T or staggered junction	2759
1	Not at junction or within 20 metres	1371

2	Crossroads	1234
3	unknown (self reported)	1087
4	Other junction	733
5	More than 4 arms (not roundabout)	292
6	Roundabout	268
7	Private drive or entrance	139
8	Slip road	84
9	Mini-roundabout	83



Accidents were rarer at slip roads and mini-roundabouts (<100), with higher rates at busy junctions like T-junctions and crossroads.

junction_control

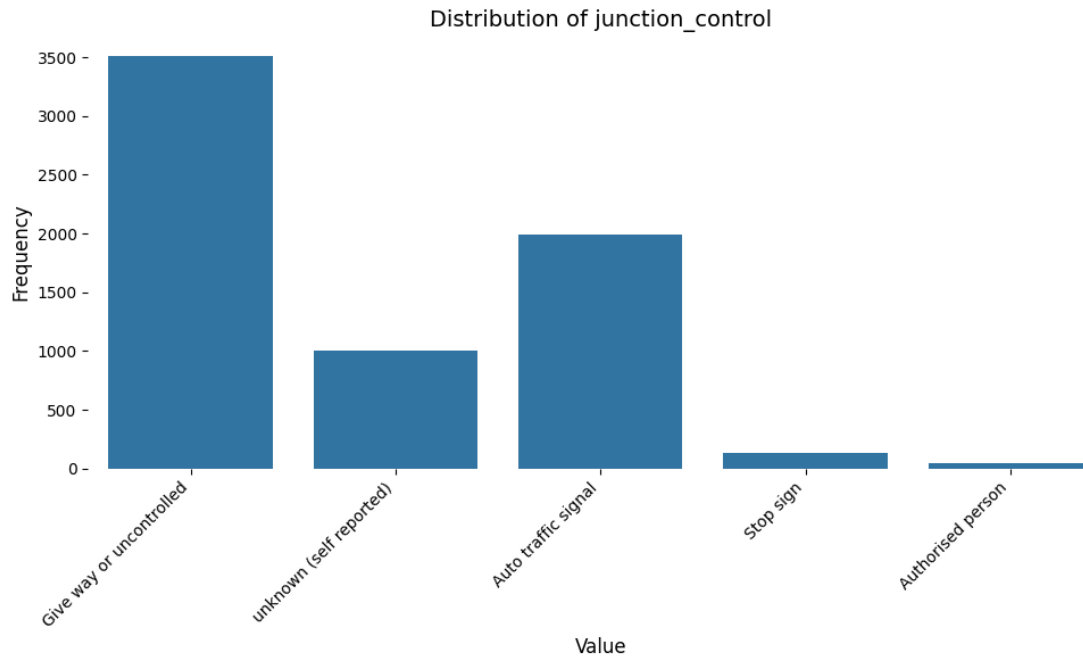
```
[77]: #Use the previously created functions to describe and plot the categorical
      ↪variable
describe_column(X_train, 'junction_control')
plot_countplot(X_train, 'junction_control')
```

Descriptive Statistics for junction_control:

Mode: Give way or uncontrolled

	Category	Frequency
0	Give way or uncontrolled	3506
1	Auto traffic signal	1991

2	unknown (self reported)	1004
3	Stop sign	136
4	Authorised person	42



Junction controls such as “Stop sign” and “Authorised person” had the lowest number of accidents, while give-way signs saw the highest accident rates.

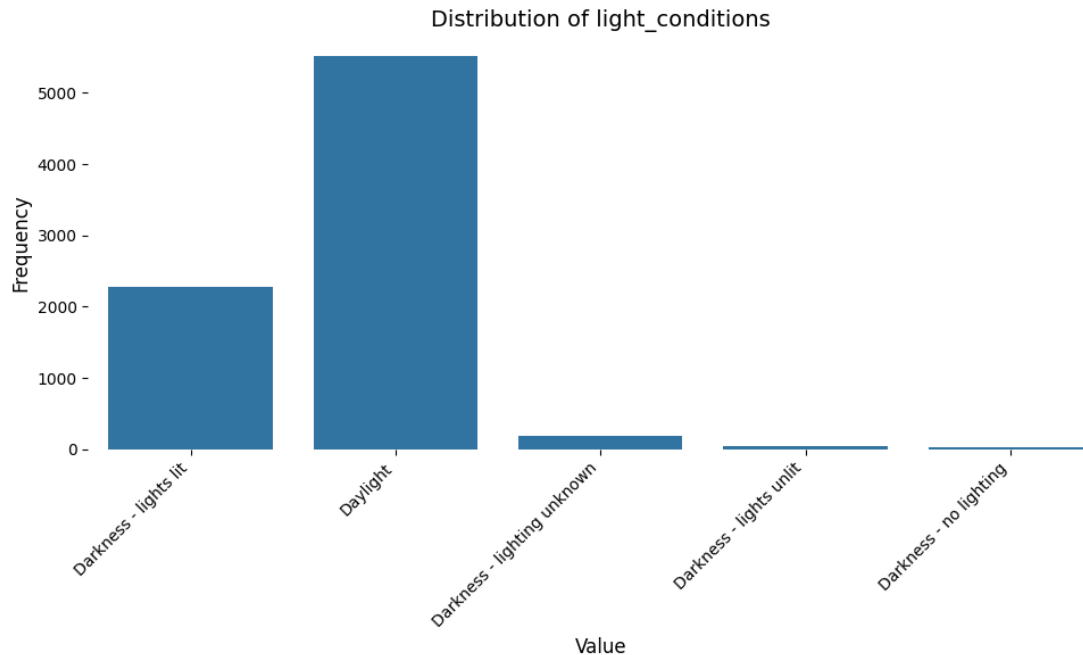
light_conditions

```
[80]: #Use the previously created functions to describe and plot the categorical
      ↪variable
      describe_column(X_train, 'light_conditions')
      plot_countplot(X_train, 'light_conditions')
```

Descriptive Statistics for light_conditions:

Mode: Daylight

	Category	Frequency
0	Daylight	5523
1	Darkness - lights lit	2273
2	Darkness - lighting unknown	180
3	Darkness - lights unlit	40
4	Darkness - no lighting	34



Accidents were the most common in daylight conditions (5523), while poorly lit or unlit conditions saw far fewer, highlighting the importance of visibility.


weather_conditions

```
[83]: #Use the previously created functions to describe and plot the categorical
      ↪variable
      describe_column(X_train, 'weather_conditions')
      plot_countplot(X_train, 'weather_conditions')
```

Descriptive Statistics for weather_conditions:

Mode: Fine no high winds

	Category	Frequency
0	Fine no high winds	5952
1	Raining no high winds	854
2	Unknown	766
3	Other	364
4	Raining + high winds	48
5	Fine + high winds	41
6	Fog or mist	13
7	Snowing no high winds	12



Most accidents occurred in Fine conditions with no high winds (5952), while the four weather conditions starting from right saw <50 each.

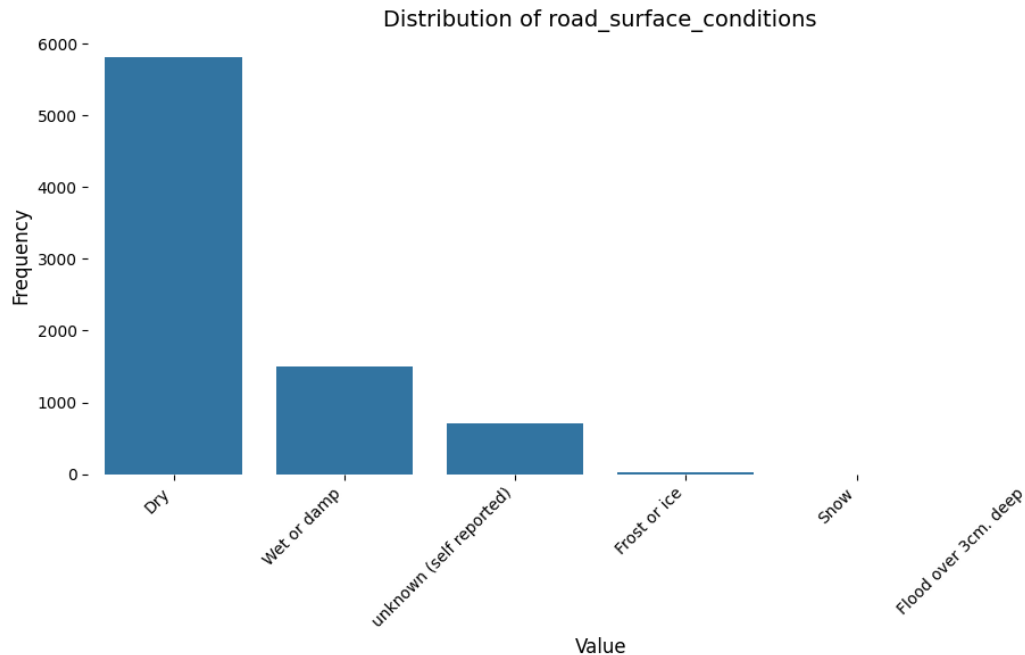
road_surface_conditions

```
[86]: #Use the previously created functions to describe and plot the categorical
      ↪variable
      describe_column(X_train, 'road_surface_conditions')
      plot_countplot(X_train, 'road_surface_conditions')
```

Descriptive Statistics for road_surface_conditions:

Mode: Dry

	Category	Frequency
0	Dry	5817
1	Wet or damp	1503
2	unknown (self reported)	705
3	Frost or ice	22
4	Flood over 3cm. deep	2
5	Snow	1



Most accidents occurred on dry roads (5,817), followed by wet conditions (1,503), with frost, ice, and flooding incidents being rare (<22).

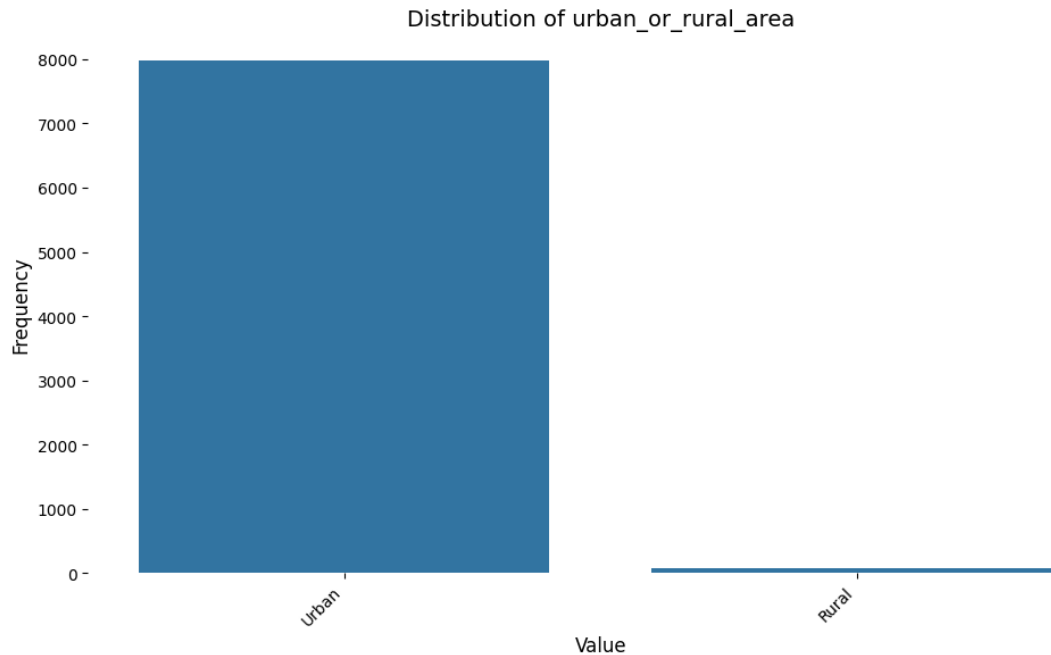
urban_or_rural_area

```
[89]: #Use the previously created functions to describe and plot the categorical
      ↪variable
      describe_column(X_train, 'urban_or_rural_area')
      plot_countplot(X_train, 'urban_or_rural_area')
```

Descriptive Statistics for urban_or_rural_area:

Mode: Urban

	Category	Frequency
0	Urban	7984
1	Rural	66



Urban areas saw 7,984 accidents vs. 66 in rural areas, highlighting higher risk in densely populated areas.

vehicle_location_restricted_lane

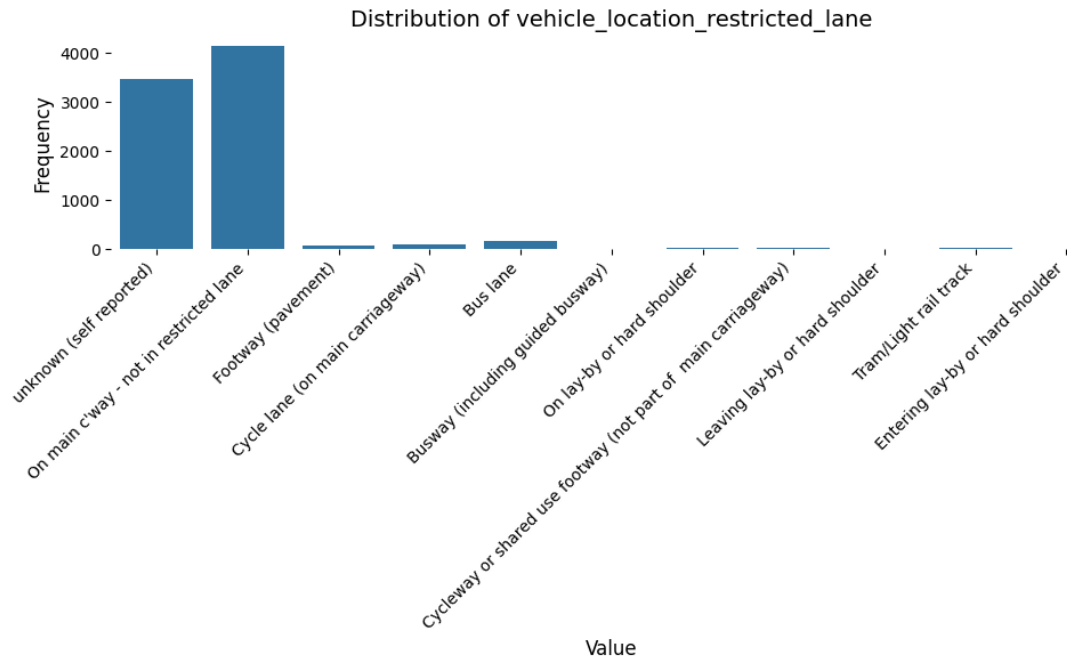
[92]: *#Use the previously created functions to describe and plot the categorical variable*

```
describe_column(X_train, 'vehicle_location_restricted_lane')
plot_countplot(X_train, 'vehicle_location_restricted_lane')
```

Descriptive Statistics for vehicle_location_restricted_lane:

Mode: On main c'way - not in restricted lane

	Category	Frequency
0	On main c'way - not in restricted lane	4140
1	unknown (self reported)	3479
2	Bus lane	172
3	Cycle lane (on main carriageway)	97
4	Footway (pavement)	80
5	Cycleway or shared use footway (not part of m...	29
6	On lay-by or hard shoulder	23
7	Tram/Light rail track	13
8	Leaving lay-by or hard shoulder	7
9	Busway (including guided busway)	5
10	Entering lay-by or hard shoulder	5



Most accidents occurred on main carriageways (4,140) and unknown locations (3,479), while restricted lanes saw far fewer incidents.


junction_location

```
[95]: #Use the previously created functions to describe and plot the categorical
      ↪variable
describe_column(X_train, 'junction_location')
plot_countplot(X_train, 'junction_location')
```

Descriptive Statistics for junction_location:

Mode: unknown (self reported)

	Category	Frequency
0	unknown (self reported)	2984
1	Approaching junction or waiting/parked at junc...	1764
2	Not at or within 20 metres of junction	1371
3	Mid Junction - on roundabout or on main road	827
4	Cleared junction or waiting/parked at junction...	493
5	Entering main road	249
6	Leaving main road	221
7	Entering roundabout	73
8	Leaving roundabout	50
9	Entering from slip road	18



Out of all the reported known junction locations approaching junction or waiting/parked at junction approach was the highest (1764) followed by not at or within 20 metres of junction (1371)

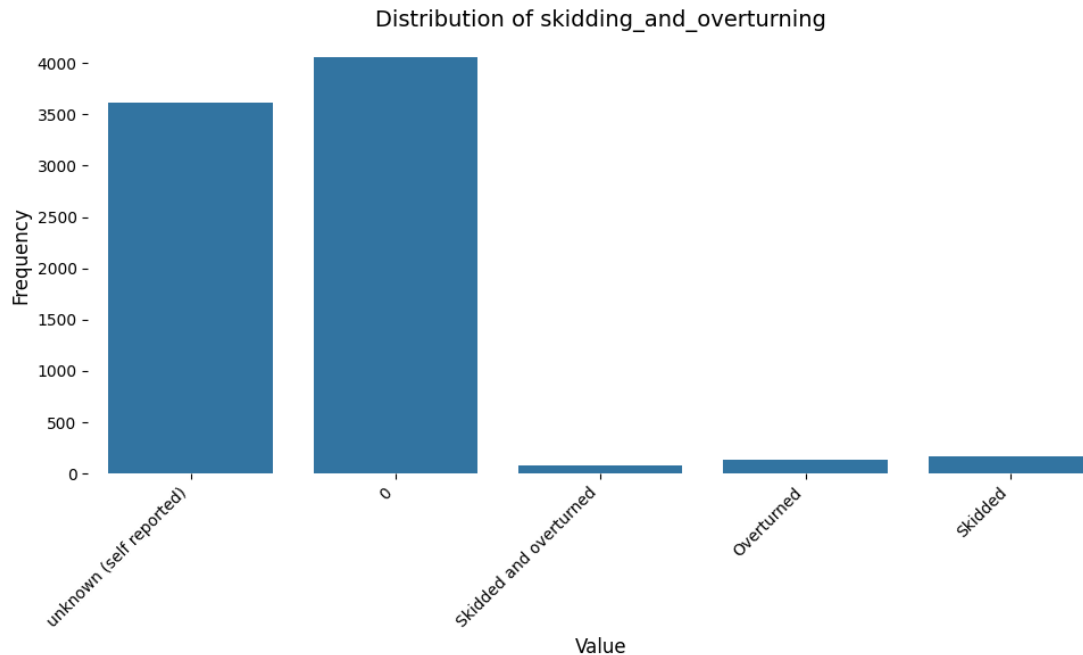
skidding_and_overturning

```
[98]: #Use the previously created functions to describe and plot the categorical
      ↪variable
describe_column(X_train, 'skidding_and_overturning')
plot_countplot(X_train, 'skidding_and_overturning')
```

Descriptive Statistics for skidding_and_overturning:

Mode: 0

	Category	Frequency
0	0	4060
1	unknown (self reported)	3613
2	Skidded	166
3	Overturned	132
4	Skidded and overturned	79



Most incidents didn't involve any skidding or overturning (4,060), but while such events were rare to happen, they were still relevant for safety improvements.

hit_object_in_carriageway

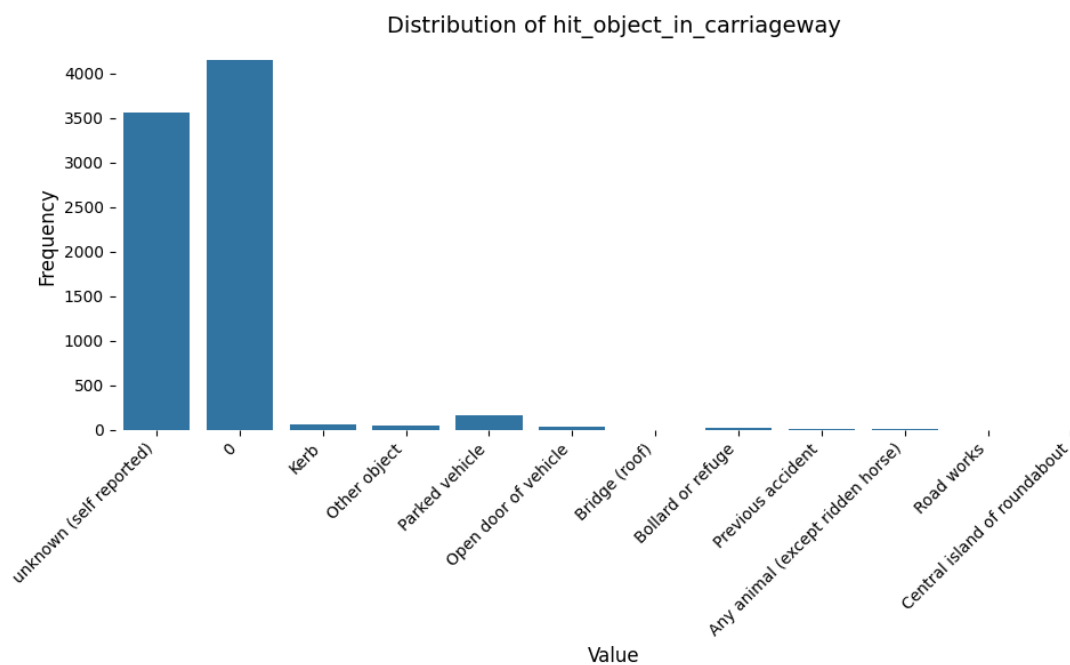
[101]: *#Use the previously created functions to describe and plot the categorical variable*

```
describe_column(X_train, 'hit_object_in_carriageway')
plot_countplot(X_train, 'hit_object_in_carriageway')
```

Descriptive Statistics for hit_object_in_carriageway:

Mode: 0

	Category	Frequency
0	0	4143
1	unknown (self reported)	3559
2	Parked vehicle	159
3	Kerb	62
4	Other object	50
5	Open door of vehicle	36
6	Bollard or refuge	28
7	Previous accident	6
8	Any animal (except ridden horse)	4
9	Bridge (roof)	1
10	Road works	1
11	Central island of roundabout	1



Most incidents didn't involve hitting objects in carriageway (4,143), with parked vehicles being the most common among the few object-related cases.

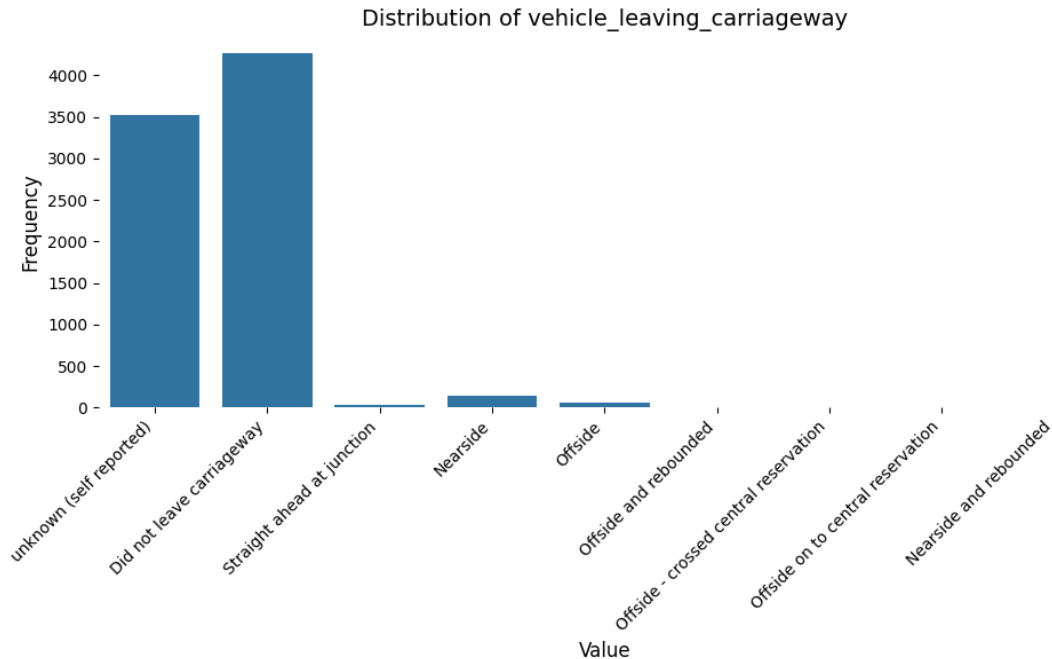
vehicle_leaving_carriageway

```
[104]: #Use the previously created functions to describe and plot the categorical
       ↪variable
describe_column(X_train, 'vehicle_leaving_carriageway')
plot_countplot(X_train, 'vehicle_leaving_carriageway')
```

Descriptive Statistics for vehicle_leaving_carriageway:

Mode: Did not leave carriageway

	Category	Frequency
0	Did not leave carriageway	4271
1	unknown (self reported)	3529
2	Nearside	147
3	Offside	56
4	Straight ahead at junction	28
5	Offside and rebounded	8
6	Nearside and rebounded	6
7	Offside on to central reservation	3
8	Offside - crossed central reservation	2



Most incidents didn't involve vehicles leaving the carriageway (4,271), but when they did, nearside departures (147) were most common followed by offside (56).

hit_object_off_carriageway


[107]: *#Use the previously created functions to describe and plot the categorical variable*

```
describe_column(X_train, 'hit_object_off_carriageway')
plot_countplot(X_train, 'hit_object_off_carriageway')
```

Descriptive Statistics for hit_object_off_carriageway:

Mode: 0

	Category	Frequency
0	0	4452
1	unknown (self reported)	3502
2	Wall or fence	27
3	Lamp post	19
4	Other permanent object	18
5	Road sign or traffic signal	13
6	Tree	6
7	Bus stop or bus shelter	6
8	Near/Offside crash barrier	4
9	Telegraph or electricity pole	2
10	Central crash barrier	1



Most incidents didn't involve hitting objects off the carriageway, but walls or fences (27), and lamp posts (19) were the most commonly struck when they did.

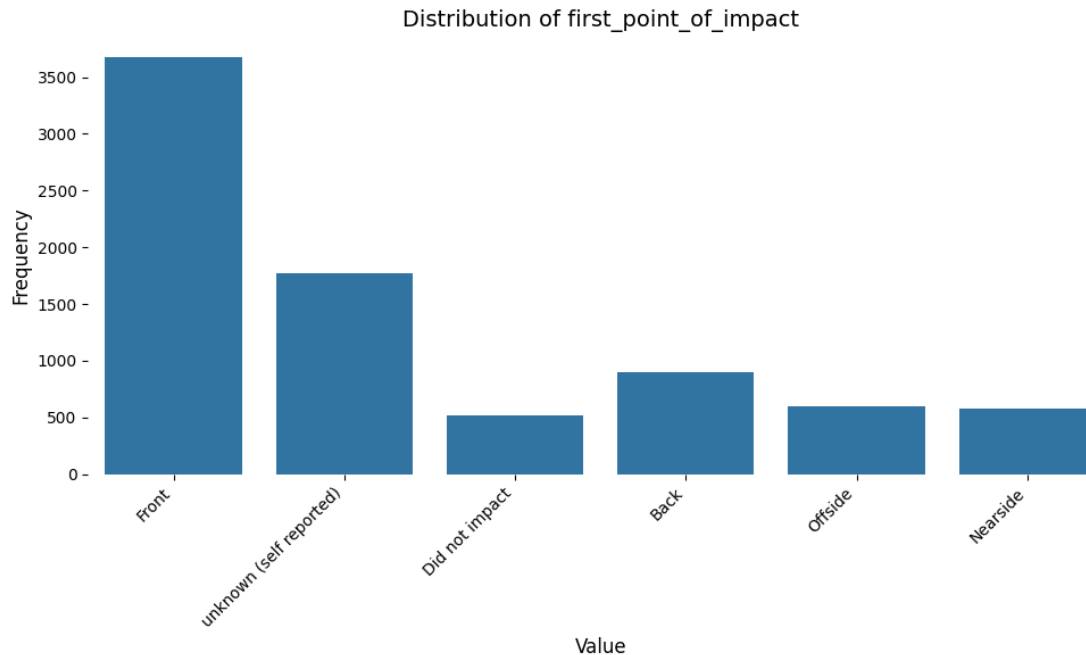
first_point_of_impact

```
[110]: #Use the previously created functions to describe and plot the categorical
       ↪variable
describe_column(X_train, 'first_point_of_impact')
plot_countplot(X_train, 'first_point_of_impact')
```

Descriptive Statistics for first_point_of_impact:

Mode: Front

	Category	Frequency
0	Front	3680
1	unknown (self reported)	1769
2	Back	904
3	Offside	600
4	Nearside	579
5	Did not impact	518



Frontal impacts were the most common (3,680), followed by rear impacts (904) amongst the reported areas of impact, highlighting the prevalence of head-on or rear-end collisions.

journey_purpose_of_driver

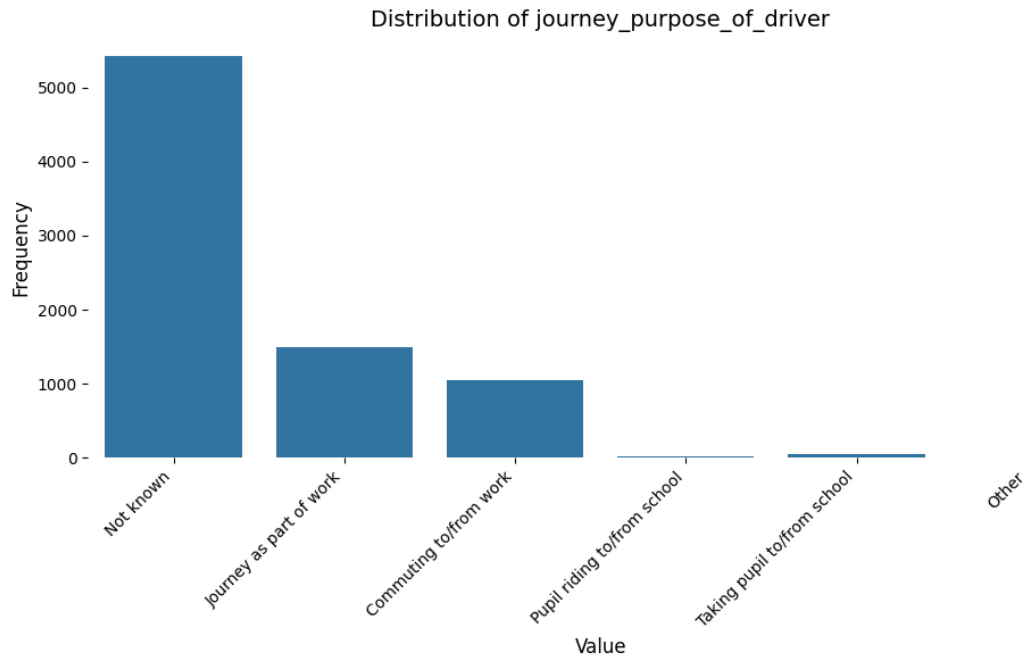
[113]: *#Use the previously created functions to describe and plot the categorical variable*

```
describe_column(X_train, 'journey_purpose_of_driver')
plot_countplot(X_train, 'journey_purpose_of_driver')
```

Descriptive Statistics for journey_purpose_of_driver:

Mode: Not known

	Category	Frequency
0	Not known	5423
1	Journey as part of work	1501
2	Commuting to/from work	1052
3	Taking pupil to/from school	57
4	Pupil riding to/from school	16
5	Other	1



Out of the known journey purpose of the driver journey as part of work involved the highest number of accidents (1501), followed by work-related travel (1,052) and commuting (1,052), with school related journeys being rare.

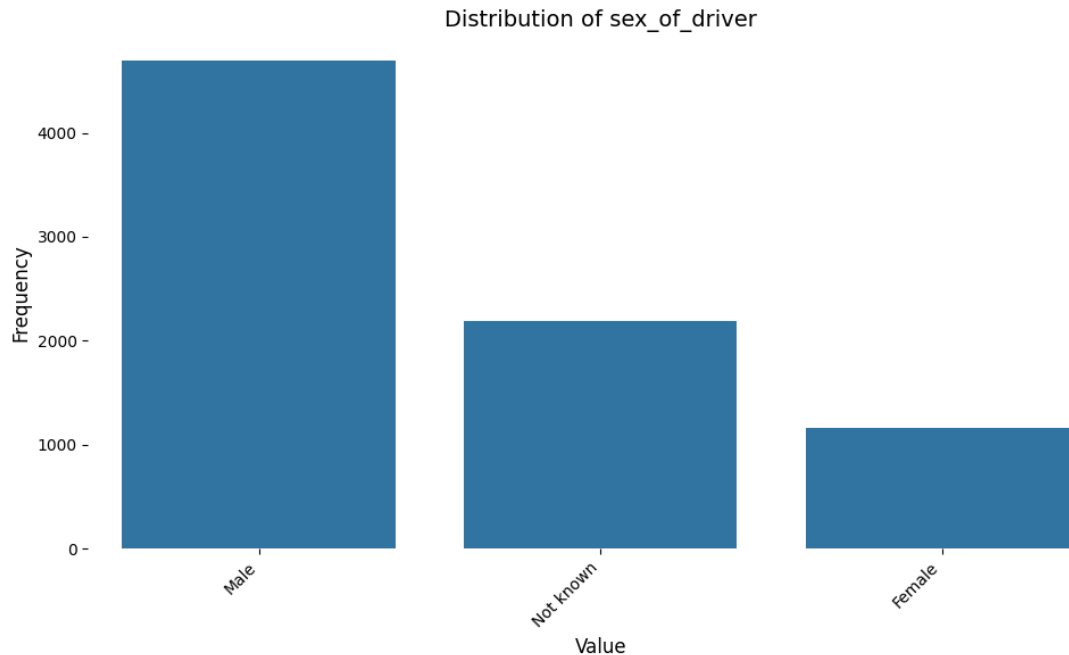
sex_of_driver

```
[116]: #Use the previously created functions to describe and plot the categorical
       ↪variable
describe_column(X_train, 'sex_of_driver')
plot_countplot(X_train, 'sex_of_driver')
```

Descriptive Statistics for sex_of_driver:

Mode: Male

	Category	Frequency
0	Male	4697
1	Not known	2186
2	Female	1167



Male drivers were involved in most incidents (4,697), but many cases (2,186) lacked gender data, highlighting gaps in demographic reporting.

vehicle_type

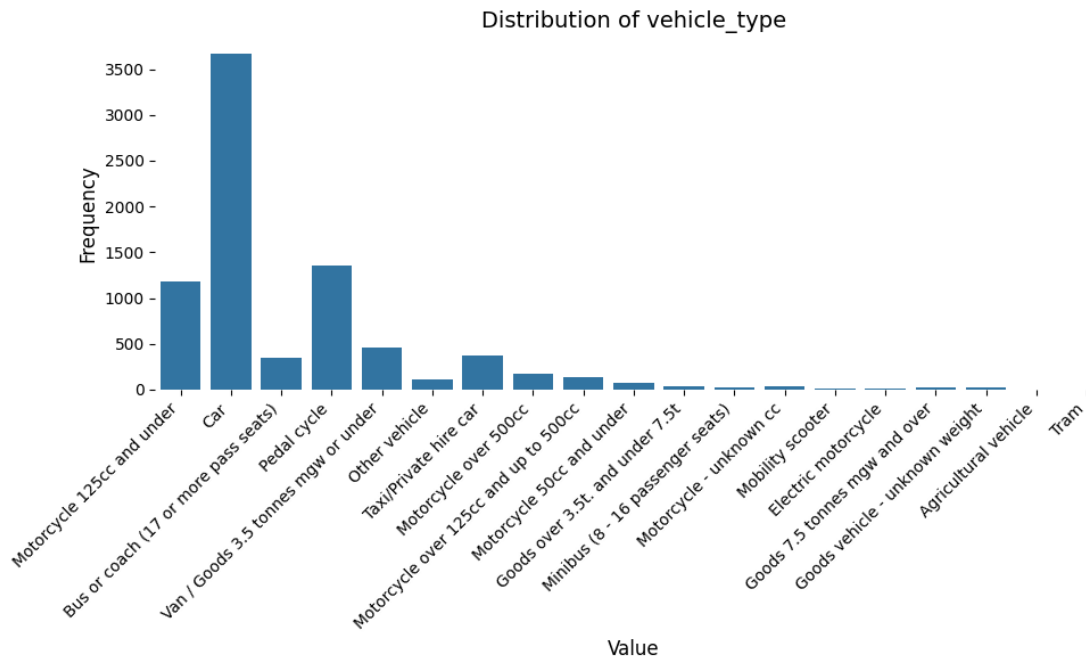
```
[119]: #Use the previously created functions to describe and plot the categorical
       ↪variable
describe_column(X_train, 'vehicle_type')
plot_countplot(X_train, 'vehicle_type')
```

Descriptive Statistics for vehicle_type:

Mode: Car

	Category	Frequency
0	Car	3667
1	Pedal cycle	1361
2	Motorcycle 125cc and under	1180
3	Van / Goods 3.5 tonnes mgw or under	454
4	Taxi/Private hire car	369
5	Bus or coach (17 or more pass seats)	352
6	Motorcycle over 500cc	180
7	Motorcycle over 125cc and up to 500cc	139
8	Other vehicle	109
9	Motorcycle 50cc and under	79
10	Motorcycle - unknown cc	35
11	Goods over 3.5t. and under 7.5t	32
12	Goods 7.5 tonnes mgw and over	25

13	Goods vehicle - unknown weight	19
14	Minibus (8 - 16 passenger seats)	19
15	Electric motorcycle	17
16	Mobility scooter	10
17	Agricultural vehicle	2
18	Tram	1



Cars were most involved in incidents (3,667), followed by pedal cycles (1,361) and small motorcycles (1,180), highlighting the dominance of personal vehicles in accidents.

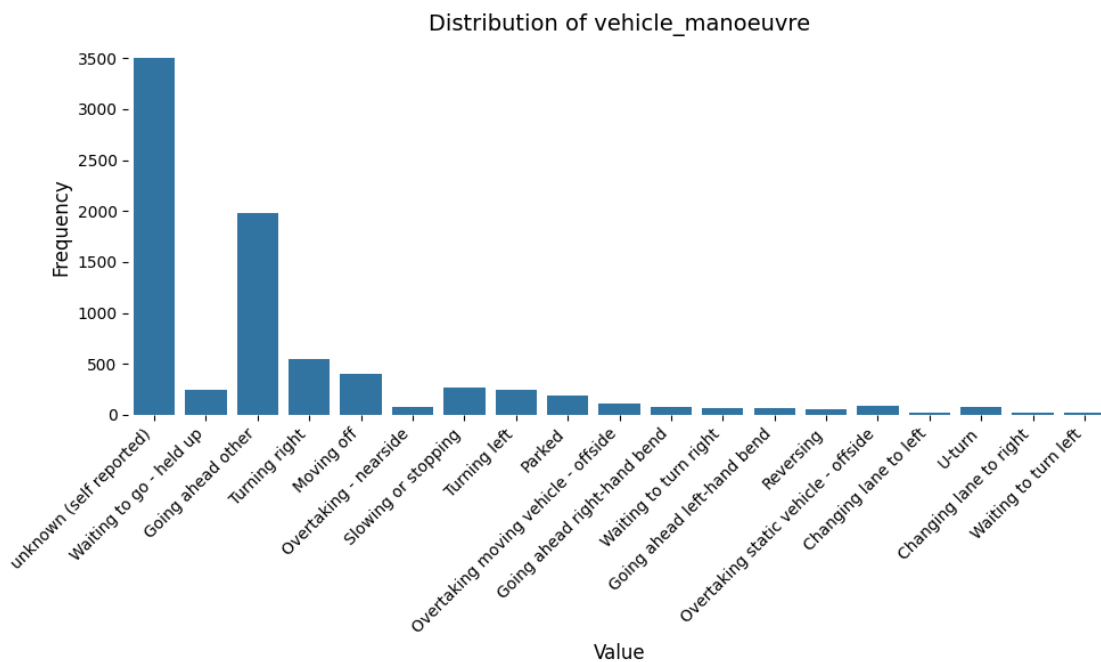
vehicle_manoeuvre

```
[122]: #Use the previously created functions to describe and plot the categorical
      ↪ variable
describe_column(X_train, 'vehicle_manoeuvre')
plot_countplot(X_train, 'vehicle_manoeuvre')
```

Descriptive Statistics for vehicle_manoeuvre:
Mode: unknown (self reported)

	Category	Frequency
0	unknown (self reported)	3503
1	Going ahead other	1985
2	Turning right	541
3	Moving off	398

4	Slowing or stopping	270
5	Turning left	248
6	Waiting to go - held up	240
7	Parked	186
8	Overtaking moving vehicle - offside	113
9	Overtaking static vehicle - offside	83
10	U-turn	81
11	Going ahead right-hand bend	79
12	Overtaking - nearside	77
13	Going ahead left-hand bend	69
14	Waiting to turn right	68
15	Reversing	49
16	Changing lane to left	24
17	Changing lane to right	19
18	Waiting to turn left	17



Most manoeuvres were marked as “unknown” (3,503), with “Going ahead other” (1,985) being the most common known action, while complex moves like lane changes were rare.

casualty_type

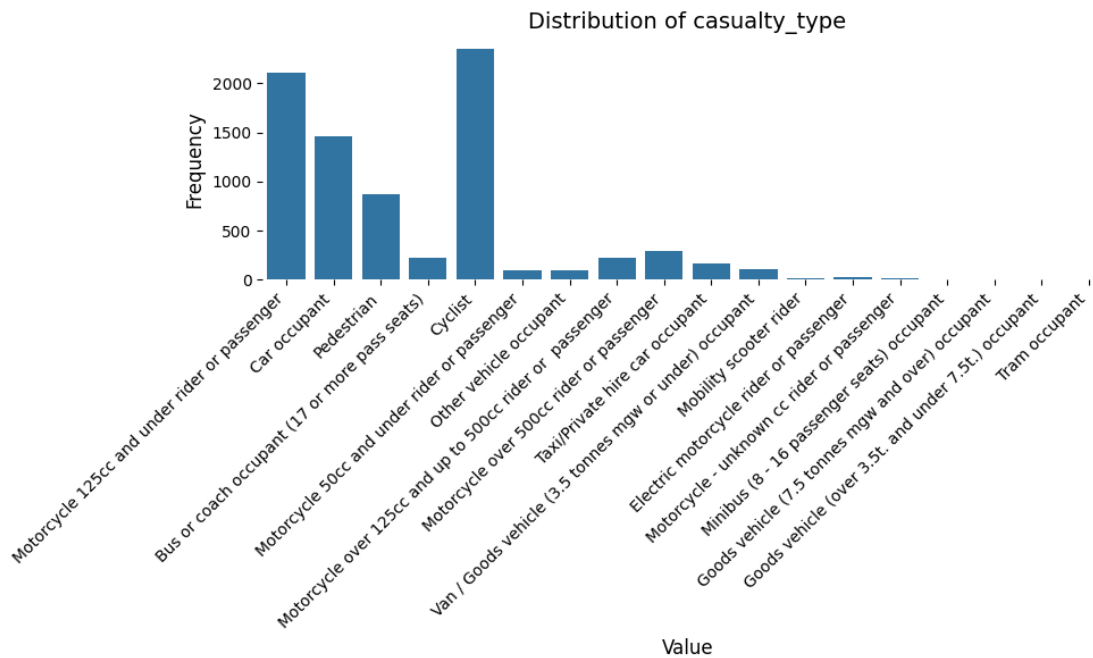
```
[125]: #Use the previously created functions to describe and plot the categorical
       ↪variable
       describe_column(X_train, 'casualty_type')
```

```
plot_countplot(X_train, 'casualty_type')
```

Descriptive Statistics for casualty_type:

Mode: Cyclist

	Category	Frequency
0	Cyclist	2344
1	Motorcycle 125cc and under rider or passenger	2104
2	Car occupant	1455
3	Pedestrian	871
4	Motorcycle over 500cc rider or passenger	290
5	Motorcycle over 125cc and up to 500cc rider or...	227
6	Bus or coach occupant (17 or more pass seats)	224
7	Taxi/Private hire car occupant	169
8	Van / Goods vehicle (3.5 tonnes mgw or under) ...	108
9	Other vehicle occupant	93
10	Motorcycle 50cc and under rider or passenger	92
11	Electric motorcycle rider or passenger	28
12	Motorcycle - unknown cc rider or passenger	17
13	Mobility scooter rider	15
14	Minibus (8 - 16 passenger seats) occupant	6
15	Goods vehicle (over 3.5t. and under 7.5t.) occ...	3
16	Goods vehicle (7.5 tonnes mgw and over) occupant	2
17	Tram occupant	2



Cyclists and small motorcycle riders were involved in the most accidents, highlighting higher risks for vulnerable two-wheel road users.

5.2 Bivariate Analysis

Relationships between the target and independent variables were explored using visualizations to assess their potential impact.

```
[129]: #Create a copy y_train dataset that has all its numerical codes converted into
        ↪their actual categorical labels:
y_train_with_labels = pd.DataFrame(y_train).copy()

for column in y_train_with_labels.columns:
    if column in guide_data['field name'].values:
        # Creating a dictionary to map code/format to label for each field_name
        mapping_dict = guide_data[guide_data['field name'] == column].
        ↪set_index('code/format')['label'].to_dict()
        # Replacing the values in top_5_boroughs based on the mapping dictionary
        if mapping_dict:
            y_train_with_labels[column] = y_train_with_labels[column].
            ↪map(mapping_dict).fillna(y_train_with_labels[column])

y_train_with_labels.head()
```

```
[129]:      accident_severity
2142          Slight
1930          Slight
8757          Slight
9327          Slight
7125          Serious
```

```
[130]: #Merge the datasets
merged_data = pd.merge(X_train, y_train_with_labels, left_index=True,
        ↪right_index=True)


#Plot accident_severity and location_authority_ons_district on a horizontal bar
        ↪chart
ax = sns.countplot(data=merged_data, x='local_authority_ons_district',
        ↪hue='accident_severity', palette='Reds')

#Plot the counts of the bars
for container in ax.containers:
    ax.bar_label(container, fontsize=10)

#Move the legend to the right of the plot and center the plot
plt.title('Accident Severity Distribution within Local Authorities')
```

```
plt.legend(title='Accident Severity', bbox_to_anchor=(1.05, 0.5), loc='center_↵left')
plt.xlabel('Local Authorities')
plt.ylabel('Count')

#Cleaner borders
sns.despine()
plt.show()
```



Across five London boroughs, slight accidents dominate, with serious cases less common and fatal ones rare.

```
[132]: #Line graph for the trend of severity on accidents in a 24 hour period
#Extract the hour from 'time'
merged_data['hour_only'] = merged_data['time'].dt.hour

#Grouping by 'hour_only' and 'accident_severity' to get counts
severity_counts = merged_data.groupby(['hour_only', 'accident_severity']).
    size().unstack(fill_value=0)

#Adding a total column for the total number of accidents in each hour
severity_counts['Total'] = severity_counts.sum(axis=1)

#Plotting the trend for each accident severity level and total accidents
plt.figure(figsize=(10, 6))
```

```

#Loop through each unique accident severity level and plot it
for severity in severity_counts.columns[:-1]: #Exclude 'Total' column from
    ↳this loop
    plt.plot(severity_counts.index, severity_counts[severity], marker='o',
    ↳label=f"Severity {severity}")

#Plotting the total number of accidents line
plt.plot(severity_counts.index, severity_counts['Total'], marker='o',
    ↳label="Total Accidents", linestyle='--', color='black')


#Adding labels for each point on the total line
for x, y in zip(severity_counts.index, severity_counts['Total']):
    plt.text(x, y, str(int(y)), color='black', ha='center', va='bottom')

#Adding labels and title
plt.title("Accident Severity Trend Over a 24-hour Period", fontsize=14)
plt.xlabel("Hours of the day", fontsize=12)
plt.ylabel("Number of Accidents", fontsize=12)

#Adding a legend
plt.legend(title="Accident Severity", loc="upper left")

#Show the plot
plt.tight_layout()
plt.xticks(severity_counts.index, severity_counts.index + 1) #Label x-axis
    ↳from 1 to 24
plt.show()

```



Accidents peak between 18:00–22:00, mostly slight in severity, with fatal and serious cases remaining low; overall trends form a U-shape, likely tied to rush hour traffic.

```
[134]: #Bar graph of accident severity by time of day
#Function below classifies periods in a day based on the hour
def categorize_time(time):
    #Extract hour from the timestamp
    hour = time.hour
    if 0 <= hour < 6:
        return "Night"
    elif 6 <= hour < 12:
        return "Morning"
    elif 12 <= hour < 18:
        return "Afternoon"
    elif 18 <= hour < 21:
        return "Evening"
    else:
        return "Night"

#Apply categorization to 'time' column
merged_data['Time of Day'] = merged_data['time'].apply(categorize_time)

#Group by 'Time of Day' and 'Accident Severity' to get counts
severity_counts = merged_data.groupby(['Time of Day', 'accident_severity']).
    .size().unstack(fill_value=0)
```

```

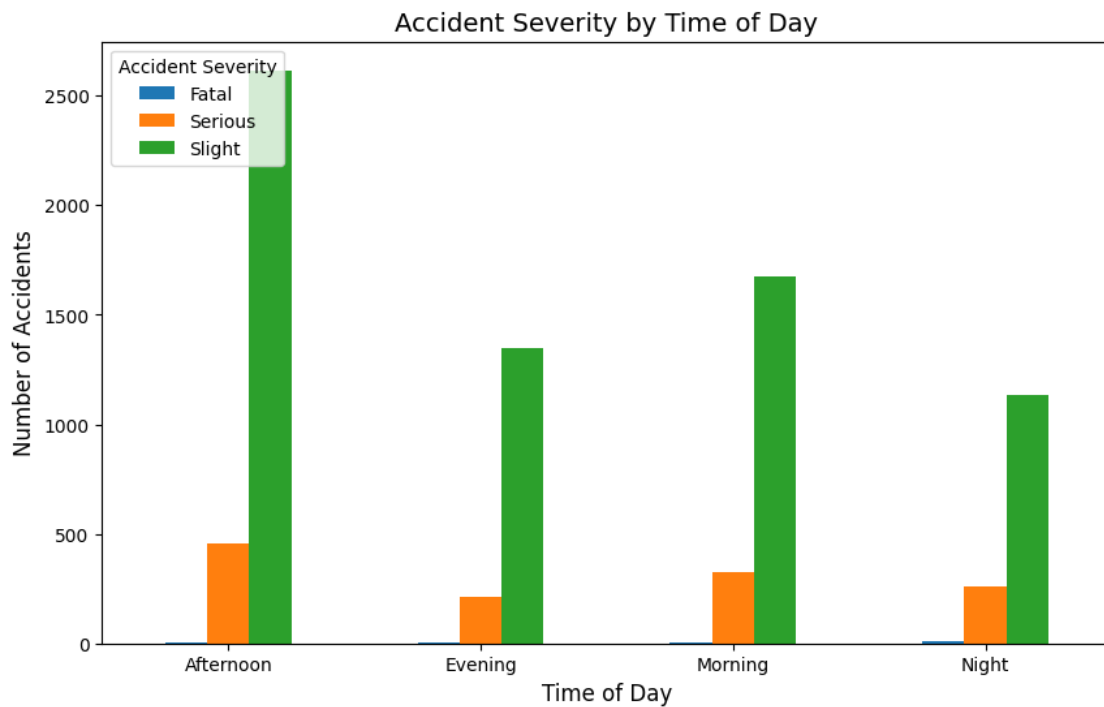
# Plotting the bar graph
ax= severity_counts.plot(kind='bar', figsize=(10, 6))
ax.set_xticklabels(ax.get_xticklabels(), rotation=0, ha='center')

#Adding labels
for line in ax.lines:
    for x, y in zip(line.get_xdata(), line.get_ydata()):
        ax.text(x, y, str(int(y)), color='black', ha='center', va='bottom')

#Adding labels and title
plt.title("Accident Severity by Time of Day", fontsize=14)
plt.xlabel("Time of Day", fontsize=12)
plt.ylabel("Number of Accidents", fontsize=12)
plt.legend(title='Accident Severity', loc="upper left")

#Show the plot
plt.show()

```



Most accidents occur in the afternoon (12–18h), with slight accidents dominating; mornings see slightly more incidents than evenings.

[136]: `#Plot accident_severity and sex_of_driver on a horizontal bar chart`


```


ax = sns.countplot(data=merged_data, x='sex_of_driver',
    hue='accident_severity', palette='Reds')

#Plot the counts of the bars
for container in ax.containers:
    ax.bar_label(container, fontsize=10)

#Move the legend to the right of the plot and center the plot
#Label the axes
plt.title('Accident Severity Distribution Across Sexes of the Drivers')
plt.legend(title='Accident Severity', bbox_to_anchor=(1.05, 0.5), loc='center_
    left')
plt.xlabel('Sex of Drivers')
plt.ylabel('Count')

#Cleaner borders
sns.despine()
plt.show()

```



Males are involved in more accidents across all severities, with 1,973 slight, 854 serious, and 17 fatal cases, while females show lower figures.

```

[138]: #Box plot of the accident severity and ages of the drivers
sns.boxplot(data=merged_data, x='accident_severity', y='age_of_driver',
    palette='Reds')

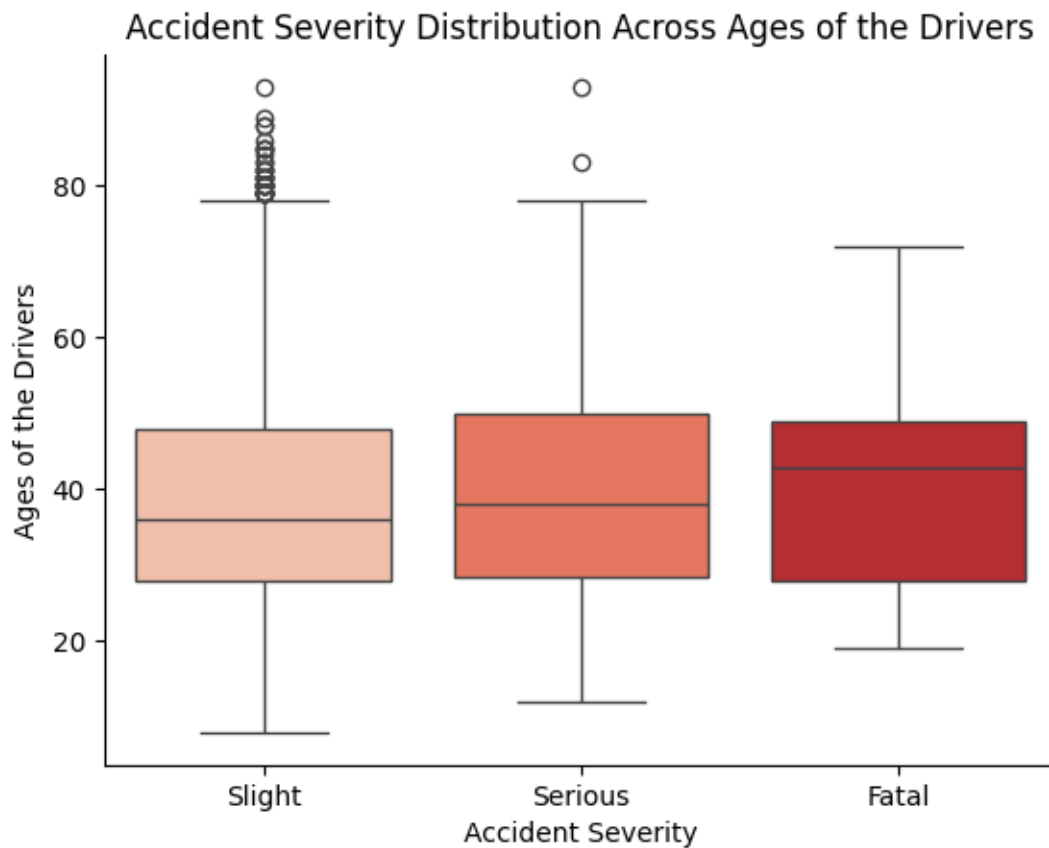
```

```

#Move the legend to the right of the plot and center the plot
plt.title('Accident Severity Distribution Across Ages of the Drivers')
plt.xlabel('Accident Severity')
plt.ylabel('Ages of the Drivers')

#Cleaner borders
sns.despine()
plt.show()

```



The boxplot shows that fatal accidents are more common among older drivers, while slight accidents compared with serious and fatal involve a younger age range.

6 Data Preprocessing

6.1 Handling Missing Values

```
[142]: #Any missing values in both the training and testing datasets will be imputed,
        ↪ using Iterative Imputer.
        #To do that, dummy variables will be created from the categorical variables.

        #Create dummy variables from the categorical variables, which are listed below
cat_cols = ['local_authority_ons_district', 'road_type', 'junction_detail',
            'junction_control', 'light_conditions', 'weather_conditions',
            'road_surface_conditions', 'urban_or_rural_area',
            'vehicle_location_restricted_lane', 'junction_location',
            'skidding_and_overturning',
            'hit_object_in_carriageway', 'vehicle_leaving_carriageway',
            'hit_object_off_carriageway', 'first_point_of_impact',
            'journey_purpose_of_driver', 'sex_of_driver', 'vehicle_type',
            'vehicle_manoeuvre', 'casualty_type']

        #Fit an encoder and transform the **trainset**
cat_vals = X_train[cat_cols]
cat_vals = X_test[cat_cols]

        #Replace 0 with 'No' in the specified categorical columns
X_train[cat_cols] = X_train[cat_cols].replace(0, 'No')
X_test[cat_cols] = X_test[cat_cols].replace(0, 'No')

one_hot_encoder = OneHotEncoder(drop="first", sparse_output=False)
```

```
[143]: #Fit and transform the training data
transformed = one_hot_encoder.fit_transform(X_train[cat_cols])

        #Get the new column names after one-hot encoding
new_col_names = one_hot_encoder.get_feature_names_out(cat_cols)

        #Add the transformed columns to X_train
for i, new_col_name in enumerate(new_col_names):
    X_train[new_col_name] = transformed[:, i]

        #Check the updated X_train
X_train.head()
```

```
[143]:
```

	accident_index	time	number_of_vehicles	\
2142	2023010435768	1900-01-01 19:35:00	2.0	
1930	2023010434377	1900-01-01 15:09:00	1.0	
8757	2023010477755	1900-01-01 17:00:00	2.0	
9327	2023010481441	1900-01-01 17:15:00	2.0	
7125	2023010467051	1900-01-01 18:40:00	2.0	

	number_of_casualties	date	local_authority_ons_district	\
2142	1.0	2023-02-27	Lambeth	
1930	1.0	2023-03-24	Lambeth	
8757	1.0	2023-11-10	Lambeth	
9327	1.0	2023-12-01	Lambeth	
7125	1.0	2023-09-19	Wandsworth	

	road_type	speed_limit	junction_detail	\
2142	Unknown	30	Not at junction or within 20 metres	
1930	Unknown	20	Not at junction or within 20 metres	
8757	Single carriageway	30	T or staggered junction	
9327	Unknown	20	unknown (self reported)	
7125	Single carriageway	20	Not at junction or within 20 metres	

	junction_control	...	\
2142	NaN	...	
1930	NaN	...	
8757	Give way or uncontrolled	...	
9327	unknown (self reported)	...	
7125	NaN	...	

	casualty_type_Motorcycle - unknown cc rider or passenger	\
2142	0.0	
1930	0.0	
8757	0.0	
9327	0.0	
7125	0.0	

	casualty_type_Motorcycle 125cc and under rider or passenger	\
2142	1.0	
1930	0.0	
8757	1.0	
9327	0.0	
7125	0.0	

	casualty_type_Motorcycle 50cc and under rider or passenger	\
2142	0.0	
1930	0.0	
8757	0.0	
9327	0.0	
7125	0.0	

	casualty_type_Motorcycle over 125cc and up to 500cc rider or passenger	\
2142	0.0	
1930	0.0	
8757	0.0	

9327	0.0
7125	0.0

	casualty_type_Motorcycle over 500cc rider or passenger \
2142	0.0
1930	0.0
8757	0.0
9327	0.0
7125	0.0

	casualty_type_Other vehicle occupant	casualty_type_Pedestrian \
2142	0.0	0.0
1930	0.0	0.0
8757	0.0	0.0
9327	0.0	1.0
7125	0.0	0.0

	casualty_type_Taxi/Private hire car occupant	casualty_type_Tram occupant \
2142	0.0	0.0
1930	0.0	0.0
8757	0.0	0.0
9327	0.0	0.0
7125	0.0	0.0

	casualty_type_Van / Goods vehicle (3.5 tonnes mgw or under) occupant
2142	0.0
1930	0.0
8757	0.0
9327	0.0
7125	0.0

[5 rows x 184 columns]

```
[144]: #Do the same for the test data
#Transform the testing data
transformed = one_hot_encoder.transform(X_test[cat_cols])

#Get the new column names after one-hot encoding
new_col_names = one_hot_encoder.get_feature_names_out(cat_cols)

#Add the transformed columns to X_test
for i, new_col_name in enumerate(new_col_names):
    X_test[new_col_name] = transformed[:, i]

#Check the updated X_train
X_test.head()
```

```

[144]:      accident_index      time  number_of_vehicles  \
8631    2023010476974  1900-01-01  08:10:00          2.0
9500    2023010482861  1900-01-01  17:45:00          2.0
1958    2023010434525  1900-01-01  18:00:00          3.0
4498    2023010450998  1900-01-01  17:55:00          2.0
7743    2023010471555  1900-01-01  11:58:00          2.0

      number_of_casualties      date local_authority_ons_district  \
8631                    2.0  2023-11-08                Croydon
9500                    1.0  2023-12-06                Croydon
1958                    1.0  2023-03-27                Southwark
4498                    1.0  2023-06-23                Lambeth
7743                    1.0  2023-10-12                Lambeth

      road_type  speed_limit      junction_detail  \
8631  Single carriageway      30  Not at junction or within 20 metres
9500          Unknown      30          unknown (self reported)
1958  Single carriageway      30          T or staggered junction
4498  Single carriageway      20          T or staggered junction
7743  Single carriageway      20  More than 4 arms (not roundabout)

      junction_control  ...  \
8631                  NaN  ...
9500  unknown (self reported)  ...
1958  Give way or uncontrolled  ...
4498  Give way or uncontrolled  ...
7743    Auto traffic signal  ...

      casualty_type_Motorcycle - unknown cc rider or passenger  \
8631                                                            0.0
9500                                                            0.0
1958                                                            0.0
4498                                                            0.0
7743                                                            0.0

      casualty_type_Motorcycle 125cc and under rider or passenger  \
8631                                                            0.0
9500                                                            1.0
1958                                                            1.0
4498                                                            0.0
7743                                                            0.0

      casualty_type_Motorcycle 50cc and under rider or passenger  \
8631                                                            0.0
9500                                                            0.0
1958                                                            0.0
4498                                                            0.0

```

7743	0.0
------	-----

casualty_type_Motorcycle over 125cc and up to 500cc rider or passenger \	
8631	0.0
9500	0.0
1958	0.0
4498	0.0
7743	0.0

casualty_type_Motorcycle over 500cc rider or passenger \	
8631	0.0
9500	0.0
1958	0.0
4498	0.0
7743	0.0

casualty_type_Other vehicle occupant casualty_type_Pedestrian \		
8631	0.0	0.0
9500	0.0	0.0
1958	0.0	0.0
4498	0.0	0.0
7743	0.0	0.0

casualty_type_Taxi/Private hire car occupant casualty_type_Tram occupant \		
8631	0.0	0.0
9500	0.0	0.0
1958	0.0	0.0
4498	0.0	0.0
7743	0.0	0.0

casualty_type_Van / Goods vehicle (3.5 tonnes mgw or under) occupant	
8631	0.0
9500	0.0
1958	0.0
4498	0.0
7743	0.0

[5 rows x 184 columns]

```
[145]: #With the dummy variables created, drop the categorical variables that were
        ↪used to create them
        #List of columns to delete
        columns_to_delete =
        ↪['local_authority_ons_district','road_type','junction_detail',
          'junction_control','light_conditions','weather_conditions',
          'road_surface_conditions','urban_or_rural_area',
          'vehicle_location_restricted_lane','junction_location',
```

```
'skidding_and_overturning', 'hit_object_in_carriageway',
'vehicle_leaving_carriageway', 'hit_object_off_carriageway',
'first_point_of_impact', 'journey_purpose_of_driver',
'sex_of_driver', 'vehicle_type', 'vehicle_manoeuvre',
'casualty_type']
```

```
#Drop the columns from the DataFrame
```

```
X_train = X_train.drop(columns=columns_to_delete)
```

```
[146]: #Check if the columns are dropped
```

```
X_train.head()
```

```
[146]:
```

	accident_index	time	number_of_vehicles	\
2142	2023010435768	1900-01-01 19:35:00	2.0	
1930	2023010434377	1900-01-01 15:09:00	1.0	
8757	2023010477755	1900-01-01 17:00:00	2.0	
9327	2023010481441	1900-01-01 17:15:00	2.0	
7125	2023010467051	1900-01-01 18:40:00	2.0	

	number_of_casualties	date	speed_limit	age_of_driver	\
2142	1.0	2023-02-27	30	26.0	
1930	1.0	2023-03-24	20	NaN	
8757	1.0	2023-11-10	30	25.0	
9327	1.0	2023-12-01	20	64.0	
7125	1.0	2023-09-19	20	60.0	

	local_authority_ons_district_Lambeth	\
2142	1.0	
1930	1.0	
8757	1.0	
9327	1.0	
7125	0.0	

	local_authority_ons_district_Southwark	\
2142	0.0	
1930	0.0	
8757	0.0	
9327	0.0	
7125	0.0	

	local_authority_ons_district_Wandsworth	...	\
2142	0.0	...	
1930	0.0	...	
8757	0.0	...	
9327	0.0	...	
7125	1.0	...	

	casualty_type_Motorcycle - unknown cc rider or passenger \
2142	0.0
1930	0.0
8757	0.0
9327	0.0
7125	0.0

	casualty_type_Motorcycle 125cc and under rider or passenger \
2142	1.0
1930	0.0
8757	1.0
9327	0.0
7125	0.0

	casualty_type_Motorcycle 50cc and under rider or passenger \
2142	0.0
1930	0.0
8757	0.0
9327	0.0
7125	0.0

	casualty_type_Motorcycle over 125cc and up to 500cc rider or passenger \
2142	0.0
1930	0.0
8757	0.0
9327	0.0
7125	0.0

	casualty_type_Motorcycle over 500cc rider or passenger \
2142	0.0
1930	0.0
8757	0.0
9327	0.0
7125	0.0

	casualty_type_Other vehicle occupant	casualty_type_Pedestrian \
2142	0.0	0.0
1930	0.0	0.0
8757	0.0	0.0
9327	0.0	1.0
7125	0.0	0.0

	casualty_type_Taxi/Private hire car occupant \
2142	0.0
1930	0.0
8757	0.0
9327	0.0

7125 0.0

	casualty_type_Tram occupant \
2142	0.0
1930	0.0
8757	0.0
9327	0.0
7125	0.0

	casualty_type_Van / Goods vehicle (3.5 tonnes mgw or under) occupant
2142	0.0
1930	0.0
8757	0.0
9327	0.0
7125	0.0

[5 rows x 164 columns]

```
[147]: #List of columns to delete
columns_to_delete =_
↳['local_authority_ons_district','road_type','junction_detail',
    'junction_control','light_conditions','weather_conditions',
    'road_surface_conditions','urban_or_rural_area',
    'vehicle_location_restricted_lane','junction_location',
    'skidding_and_overturning','hit_object_in_carriageway',
    'vehicle_leaving_carriageway','hit_object_off_carriageway',
    'first_point_of_impact','journey_purpose_of_driver',
    'sex_of_driver','vehicle_type','vehicle_manoeuvre',
    'casualty_type']

#Drop the columns from the DataFrame
X_test = X_test.drop(columns=columns_to_delete)
```

```
[148]: #Check how many NaN values there are in the X_train dataset
X_train.isnull().values.sum()
```

```
[148]: np.int64(2009)
```

```
[149]: #Check how many NaN values there are in the y_train dataset
y_train.isnull().sum()
```

```
[149]: np.int64(0)
```

```
[150]: #Check how many NaN values there are in the X_test dataset
X_test.isnull().values.sum()
```

```
[150]: np.int64(483)
```

```
[151]: #Check how many NaN values there are in the y_test dataset
y_test.isnull().sum()
```

```
[151]: np.int64(0)
```

```
[152]: #Create copies of the training and testing datasets for X that only have the
      ↪ date and time columns
      #Will prevent errors from cropping up when the imputer is ran
X_train_dropped_columns = X_train[['date', 'time']]
X_test_dropped_columns = X_test[['date', 'time']]
```

```
[153]: #Date and time columns are dropped from the two X datasets
X_train = X_train.drop(columns=['date', 'time']).copy()
X_test = X_test.drop(columns=['date', 'time']).copy()
```

```
[154]: start = timer()
      #The imputer is ran on the X_train dataset
      imputer = IterativeImputer(max_iter = 30)
      X_train = pd.DataFrame(imputer.fit_transform(X_train),
                             columns = X_train.columns,
                             index = X_train.index)

      #Check to see if any NaN values are still remaining
      X_train.isnull().sum().sum()

      #Print the execution time
      print("Execution time HH:MM:SS:", timedelta(seconds=timer() - start))
```

Execution time HH:MM:SS: 0:00:34.129053

```
[155]: #Run the imputer on the X_test dataset as well
X_test = pd.DataFrame(imputer.transform(X_test),
                       columns = X_test.columns,
                       index = X_test.index)

      #Check to see if there any NaN values remaining
      X_test.isnull().sum().sum()
```

```
[155]: np.int64(0)
```

```
[156]: #Add the dropped date and time columns back into the train and test datasets
      ↪ for X
X_train = pd.concat([X_train, X_train_dropped_columns], axis=1)
X_test = pd.concat([X_test, X_test_dropped_columns], axis=1)
```

```
[157]: #Check if the changes have taken place using a visualization
pd.options.display.float_format = '{:.0f}'.format
```

```
X_train.head()
```

```
[157]:
```

	accident_index	number_of_vehicles	number_of_casualties	speed_limit	\
2142	2023010435768	2	1	30	
1930	2023010434377	1	1	20	
8757	2023010477755	2	1	30	
9327	2023010481441	2	1	20	
7125	2023010467051	2	1	20	

	age_of_driver	local_authority_ons_district_Lambeth	\
2142	26	1	
1930	40	1	
8757	25	1	
9327	64	1	
7125	60	0	

	local_authority_ons_district_Southwark	\
2142	0	
1930	0	
8757	0	
9327	0	
7125	0	

	local_authority_ons_district_Wandsworth	\
2142	0	
1930	0	
8757	0	
9327	0	
7125	1	

	local_authority_ons_district_Westminster	road_type_One way street	...	\
2142	0	0	...	
1930	0	0	...	
8757	0	0	...	
9327	0	0	...	
7125	0	0	...	

	casualty_type_Motorcycle 50cc and under rider or passenger	\
2142	0	
1930	0	
8757	0	
9327	0	
7125	0	

	casualty_type_Motorcycle over 125cc and up to 500cc rider or passenger	\
2142	0	
1930	0	

8757	0
9327	0
7125	0

	casualty_type_Motorcycle over 500cc rider or passenger \
2142	0
1930	0
8757	0
9327	0
7125	0

	casualty_type_Other vehicle occupant	casualty_type_Pedestrian \
2142	0	0
1930	0	0
8757	0	0
9327	0	1
7125	0	0

	casualty_type_Taxi/Private hire car occupant \
2142	0
1930	0
8757	0
9327	0
7125	0

	casualty_type_Tram occupant \
2142	0
1930	0
8757	0
9327	0
7125	0

	casualty_type_Van / Goods vehicle (3.5 tonnes mgw or under) occupant \
2142	0
1930	0
8757	0
9327	0
7125	0

	date	time
2142	2023-02-27 1900-01-01	19:35:00
1930	2023-03-24 1900-01-01	15:09:00
8757	2023-11-10 1900-01-01	17:00:00
9327	2023-12-01 1900-01-01	17:15:00
7125	2023-09-19 1900-01-01	18:40:00

[5 rows x 164 columns]

```
[158]: #Do the same for the X_test dataset as well
X_test.head()
```

```
[158]:      accident_index  number_of_vehicles  number_of_casualties  speed_limit  \
8631      2023010476974                2                2          30
9500      2023010482861                2                1          30
1958      2023010434525                3                1          30
4498      2023010450998                2                1          20
7743      2023010471555                2                1          20

      age_of_driver  local_authority_ons_district_Lambeth  \
8631              47                                0
9500              40                                0
1958              25                                0
4498              26                                1
7743              23                                1

      local_authority_ons_district_Southwark  \
8631                                         0
9500                                         0
1958                                         1
4498                                         0
7743                                         0

      local_authority_ons_district_Wandsworth  \
8631                                         0
9500                                         0
1958                                         0
4498                                         0
7743                                         0

      local_authority_ons_district_Westminster  road_type_One way street  ...  \
8631                                         0          0  ...
9500                                         0          0  ...
1958                                         0          0  ...
4498                                         0          0  ...
7743                                         0          0  ...

      casualty_type_Motorcycle 50cc and under rider or passenger  \
8631                          0
9500                          0
1958                          0
4498                          0
7743                          0

      casualty_type_Motorcycle over 125cc and up to 500cc rider or passenger  \
8631                          0
```

9500	0
1958	0
4498	0
7743	0

	casualty_type_Motorcycle over 500cc rider or passenger \
8631	0
9500	0
1958	0
4498	0
7743	0

	casualty_type_Other vehicle occupant	casualty_type_Pedestrian \
8631	0	0
9500	0	0
1958	0	0
4498	0	0
7743	0	0

	casualty_type_Taxi/Private hire car occupant \
8631	0
9500	0
1958	0
4498	0
7743	0

	casualty_type_Tram occupant \
8631	0
9500	0
1958	0
4498	0
7743	0

	casualty_type_Van / Goods vehicle (3.5 tonnes mgw or under) occupant \
8631	0
9500	0
1958	0
4498	0
7743	0

	date	time
8631	2023-11-08 1900-01-01	08:10:00
9500	2023-12-06 1900-01-01	17:45:00
1958	2023-03-27 1900-01-01	18:00:00
4498	2023-06-23 1900-01-01	17:55:00
7743	2023-10-12 1900-01-01	11:58:00

[5 rows x 164 columns]

6.2 Outlier Detection

To avoid bias in modeling, outliers identified during EDA will be removed using Isolation Forest.

```
[161]: #Check the size of the X_train and X_test datasets
X_train.shape, X_test.shape
```

```
[161]: ((8050, 164), (2013, 164))
```

```
[162]: #Check to see if the first few rows contain any outliers
X_train.head()
```

```
[162]:      accident_index  number_of_vehicles  number_of_casualties  speed_limit  \
2142      2023010435768                2                1          30
1930      2023010434377                1                1          20
8757      2023010477755                2                1          30
9327      2023010481441                2                1          20
7125      2023010467051                2                1          20
```

```
      age_of_driver  local_authority_ons_district_Lambeth  \
2142             26                1
1930             40                1
8757             25                1
9327             64                1
7125             60                0
```

```
      local_authority_ons_district_Southwark  \
2142                0
1930                0
8757                0
9327                0
7125                0
```

```
      local_authority_ons_district_Wandsworth  \
2142                0
1930                0
8757                0
9327                0
7125                1
```

```
      local_authority_ons_district_Westminster  road_type_One way street  ...  \
2142                0                0  ...
1930                0                0  ...
8757                0                0  ...
9327                0                0  ...
7125                0                0  ...
```


	casualty_type_Motorcycle 50cc and under rider or passenger \	
2142	0	
1930	0	
8757	0	
9327	0	
7125	0	
	casualty_type_Motorcycle over 125cc and up to 500cc rider or passenger \	
2142	0	
1930	0	
8757	0	
9327	0	
7125	0	
	casualty_type_Motorcycle over 500cc rider or passenger \	
2142	0	
1930	0	
8757	0	
9327	0	
7125	0	
	casualty_type_Other vehicle occupant casualty_type_Pedestrian \	
2142	0	0
1930	0	0
8757	0	0
9327	0	1
7125	0	0
	casualty_type_Taxi/Private hire car occupant \	
2142	0	
1930	0	
8757	0	
9327	0	
7125	0	
	casualty_type_Tram occupant \	
2142	0	
1930	0	
8757	0	
9327	0	
7125	0	
	casualty_type_Van / Goods vehicle (3.5 tonnes mgw or under) occupant \	
2142	0	
1930	0	
8757	0	

```
9327 0
7125 0
```

```
      date      time
2142 2023-02-27 1900-01-01 19:35:00
1930 2023-03-24 1900-01-01 15:09:00
8757 2023-11-10 1900-01-01 17:00:00
9327 2023-12-01 1900-01-01 17:15:00
7125 2023-09-19 1900-01-01 18:40:00
```

```
[5 rows x 164 columns]
```

```
[163]: #For the Isolation Forest to work, the date and time columns must be split into
      →year, month, day, hour, and minute.
      #Split the date and time columns in the X_train dataset into year, month, day,
      →hour, and minute columns
X_train['year'] = X_train['date'].dt.year
X_train['month'] = X_train['date'].dt.month
X_train['day'] = X_train['date'].dt.day
X_train['hour'] = X_train['time'].dt.hour
X_train['minute'] = X_train['time'].dt.minute

#Drop the time and date columns as they are no longer necessary
X_train = X_train.drop(columns=['time', 'date'])
```

```
[164]: #The Isolation Forest is trained on the X_train dataset
clf = IsolationForest(n_estimators=100, random_state=7, contamination=0.03).
      →fit(X_train)
yhat = clf.predict(X_train)
yhat
```

```
[164]: array([1, 1, 1, ..., 1, 1, 1])
```

```
[165]: start = timer()

      #Run the Isolation Forest on the X_train
X_train = X_train[yhat != -1]
y_train = y_train[yhat != -1]
X_train.shape

      #Print the execution time
print("Execution time HH:MM:SS:", timedelta(seconds=timer() - start))
```

```
Execution time HH:MM:SS: 0:00:00.004325
```

```
[166]: start = timer()

      #Do the same for the X_test
```

```

X_test['year'] = X_test['date'].dt.year
X_test['month'] = X_test['date'].dt.month
X_test['day'] = X_test['date'].dt.day
X_test['hour'] = X_test['time'].dt.hour
X_test['minute'] = X_test['time'].dt.minute
X_test = X_test.drop(columns=['date', 'time'])

hat = clf.predict(X_test)
X_test = X_test[hat == 1]
y_test = y_test[hat != -1]
X_test.shape

#Print the execution time
print("Execution time HH:MM:SS:", timedelta(seconds=timer() - start))

```

Execution time HH:MM:SS: 0:00:00.031387

6.3 Log Transformation

```


[168]: #Apply log transformation on the numerical variables so that they are in log_
      ↪ form and scaled to one another
      #Create histograms for the variables that log transformation will be conducted_
      ↪ on
      X_train[['number_of_vehicles', 'number_of_casualties', 'speed_limit',
      ↪ 'age_of_driver']].hist(bins=50, figsize=(8,8))

```

```

[168]: array([[<Axes: title={'center': 'number_of_vehicles'}>,
      <Axes: title={'center': 'number_of_casualties'}>],
      [<Axes: title={'center': 'speed_limit'}>,
      <Axes: title={'center': 'age_of_driver'}>]], dtype=object)

```

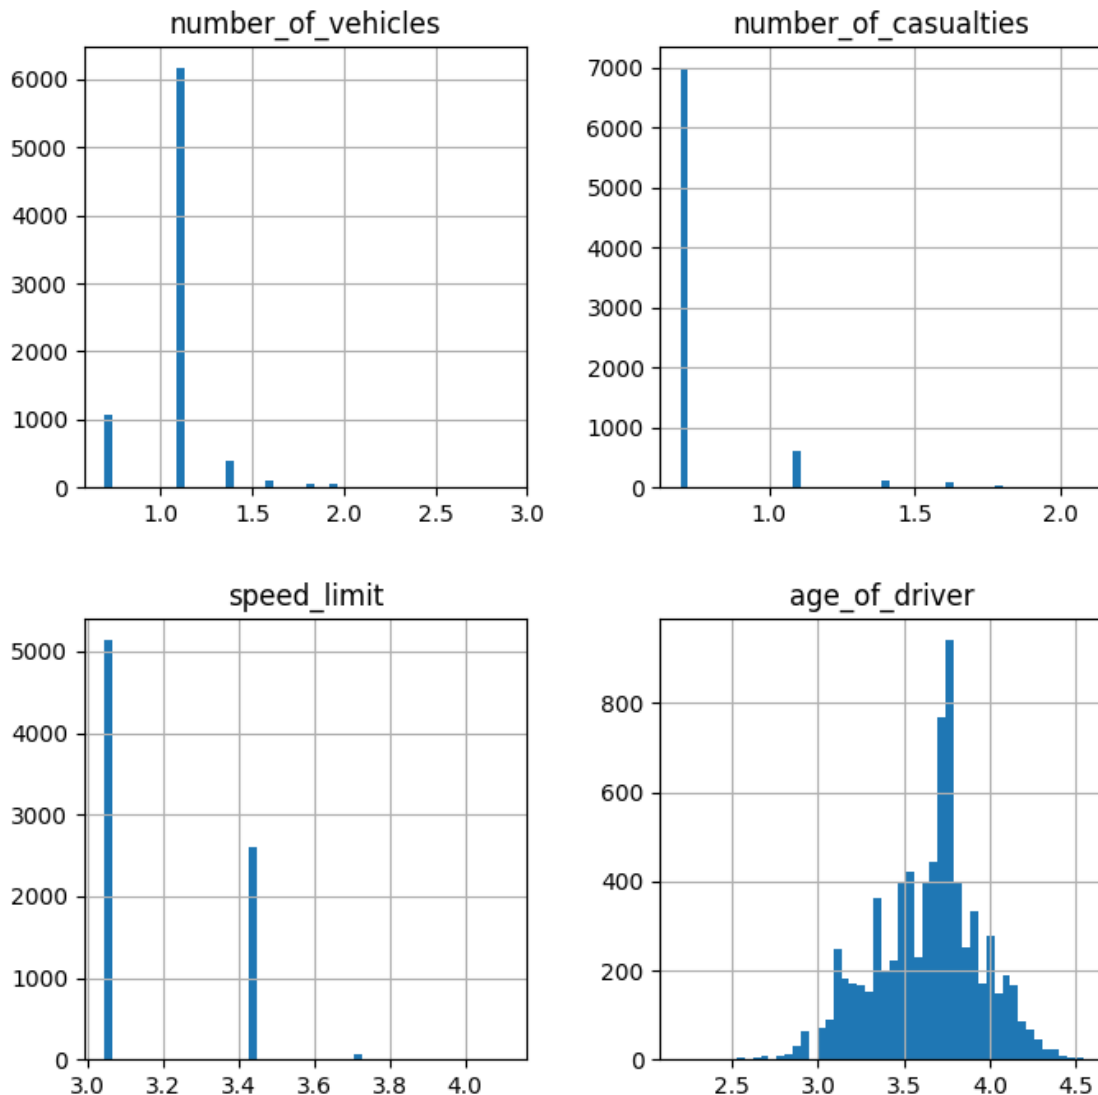


```
[169]: #Apply log transformation on the variables
for col in ['number_of_vehicles', 'number_of_casualties', 'speed_limit',
            'age_of_driver']:
    X_train.loc[:, col] = np.log(X_train.loc[:, col] + 1)
    X_test.loc[:, col] = np.log(X_test.loc[:, col] + 1)
```

```
[170]: #Create histograms again for the newly transformed variables
X_train[['number_of_vehicles', 'number_of_casualties', 'speed_limit',
        'age_of_driver']].hist(bins=50, figsize=(8,8))
```

```
[170]: array([[<Axes: title={'center': 'number_of_vehicles'}>,
               <Axes: title={'center': 'number_of_casualties'}>],
               [<Axes: title={'center': 'speed_limit'}>,
               <Axes: title={'center': 'age_of_driver'}>]])
```

```
<Axes: title=[('center': 'age_of_driver')>]], dtype=object)
```



After log transformation, the distributions of skewed variables like number of vehicles, casualties, speed limit, and driver age appear more normalized and less concentrated at lower values.

6.4 Feature Engineering

A new feature was created using year, month, and day. The new feature, day of week, will be used in the individual assignment to highlight on what day in a week an accident occurred.

```
[174]: #Create a new feature using the year, month, and day columns
X_train['date_combined'] = pd.to_datetime(X_train[['year', 'month', 'day']])

#Extract the day of the week as a number (1=Monday, 2=Tuesday, etc.)
```

```

X_train['day_of_week'] = X_train['date_combined'].dt.weekday + 1 # +1 to make_
↳ Monday = 1, Sunday = 7

#Drop the intermediate 'date_combined' column
X_train = X_train.drop(columns=['date_combined', 'year', 'day', 'minute'])

X_train.head()

```

```

[174]:
    accident_index  number_of_vehicles  number_of_casualties  speed_limit \
2142    2023010435768                1                    1          3
1930    2023010434377                1                    1          3
8757    2023010477755                1                    1          3
9327    2023010481441                1                    1          3
7125    2023010467051                1                    1          3

    age_of_driver  local_authority_ons_district_Lambeth \
2142            3                                1
1930            4                                1
8757            3                                1
9327            4                                1
7125            4                                0

    local_authority_ons_district_Southwark \
2142            0
1930            0
8757            0
9327            0
7125            0

    local_authority_ons_district_Wandsworth \
2142            0
1930            0
8757            0
9327            0
7125            1

    local_authority_ons_district_Westminster  road_type_One way street  ... \
2142            0                    0  ...
1930            0                    0  ...
8757            0                    0  ...
9327            0                    0  ...
7125            0                    0  ...

    casualty_type_Motorcycle over 125cc and up to 500cc rider or  passenger \
2142            0
1930            0
8757            0

```

9327	0
7125	0

	casualty_type_Motorcycle over 500cc rider or passenger \
2142	0
1930	0
8757	0
9327	0
7125	0

	casualty_type_Other vehicle occupant	casualty_type_Pedestrian \
2142	0	0
1930	0	0
8757	0	0
9327	0	1
7125	0	0

	casualty_type_Taxi/Private hire car occupant \
2142	0
1930	0
8757	0
9327	0
7125	0

	casualty_type_Tram occupant \
2142	0
1930	0
8757	0
9327	0
7125	0

	casualty_type_Van / Goods vehicle (3.5 tonnes mgw or under) occupant \
2142	0
1930	0
8757	0
9327	0
7125	0

	month	hour	day_of_week
2142	2	19	1
1930	3	15	5
8757	11	17	5
9327	12	17	5
7125	9	18	2

[5 rows x 165 columns]

```
[175]: #Do the same for X_test
X_test['date_combined'] = pd.to_datetime(X_test[['year', 'month', 'day']])

#Extract the day of the week as a number (1=Monday, 2=Tuesday, etc.)
X_test['day_of_week'] = X_test['date_combined'].dt.weekday + 1 # +1 to make
↳ Monday = 1, Sunday = 7

#Drop the intermediate 'date_combined' column
X_test = X_test.drop(columns=['date_combined', 'year', 'day', 'minute'])

X_test.head()
```

```
[175]:
```

	accident_index	number_of_vehicles	number_of_casualties	speed_limit	\
9500	2023010482861	1	1	3	
1958	2023010434525	1	1	3	
4498	2023010450998	1	1	3	
7743	2023010471555	1	1	3	
7059	2023010466701	1	1	3	

	age_of_driver	local_authority_ons_district_Lambeth	\
9500	4	0	
1958	3	0	
4498	3	1	
7743	3	1	
7059	3	0	

	local_authority_ons_district_Southwark	\
9500	0	
1958	1	
4498	0	
7743	0	
7059	0	

	local_authority_ons_district_Wandsworth	\
9500	0	
1958	0	
4498	0	
7743	0	
7059	0	

	local_authority_ons_district_Westminster	road_type_One way street	...	\
9500	0	0	...	
1958	0	0	...	
4498	0	0	...	
7743	0	0	...	
7059	0	0	...	

	casualty_type_Motorcycle over 125cc and up to 500cc rider or passenger \	
9500	0	
1958	0	
4498	0	
7743	0	
7059	0	

	casualty_type_Motorcycle over 500cc rider or passenger \	
9500	0	
1958	0	
4498	0	
7743	0	
7059	0	

	casualty_type_Other vehicle occupant		casualty_type_Pedestrian \	
9500	0		0	
1958	0		0	
4498	0		0	
7743	0		0	
7059	0		0	

	casualty_type_Taxi/Private hire car occupant \	
9500	0	
1958	0	
4498	0	
7743	0	
7059	0	

	casualty_type_Tram occupant \	
9500	0	
1958	0	
4498	0	
7743	0	
7059	0	

	casualty_type_Van / Goods vehicle (3.5 tonnes mgw or under) occupant \	
9500	0	
1958	0	
4498	0	
7743	0	
7059	0	

	month	hour	day_of_week
9500	12	17	3
1958	3	18	1
4498	6	17	5
7743	10	11	4

7059 9 8 1

[5 rows x 165 columns]

The fatal class and the serious class in the target variable were combined, because there were too little counts for both of the classes, as shown previously in the bar plot in the exploratory data analysis.

```
[177]: #For both of the datasets that contained the target variable, the fatal class  
       ↪and the serious class were combined  
y_train = y_train.replace(1, 2)  
y_test = y_test.replace(1, 2)
```

```
[178]: #Descriptive Statistics of the Target Variable  
  
#Frequency counts  
frequency_counts = pd.DataFrame(y_train.value_counts().sort_index())  
  
#Mode  
mode_value = y_train.mode()[0]  
  
#Median  
median_value = y_train.median()  
  
#Summary  
print("Frequency Counts:\n", frequency_counts)  
print("\nMode:", mode_value)  
print("Median:", median_value)
```

Frequency Counts:

	count
accident_severity	
2	1232
3	6576

Mode: 3

Median: 3.0

7 Conclusion

In the assignment, a new dataset was generated by picking out variables that the target variable was most related to. The actual predictive model will be created in the individual assignment; it will predict the severity of a car collision accident occurring in the London boroughs based on the independent variables. Knowing what elements influence an accident to be more severe will enable the London councils to better focus on more important areas.

8 Data Exporting

The code below exports the transformed and split datasets in this notebook to CSV files so that they can be used for the individual assignment.

```
[183]: X_train.to_csv("X_train.csv")
      y_train.to_csv("y_train.csv")
      X_test.to_csv("X_test.csv")
      y_test.to_csv("y_test.csv")
```